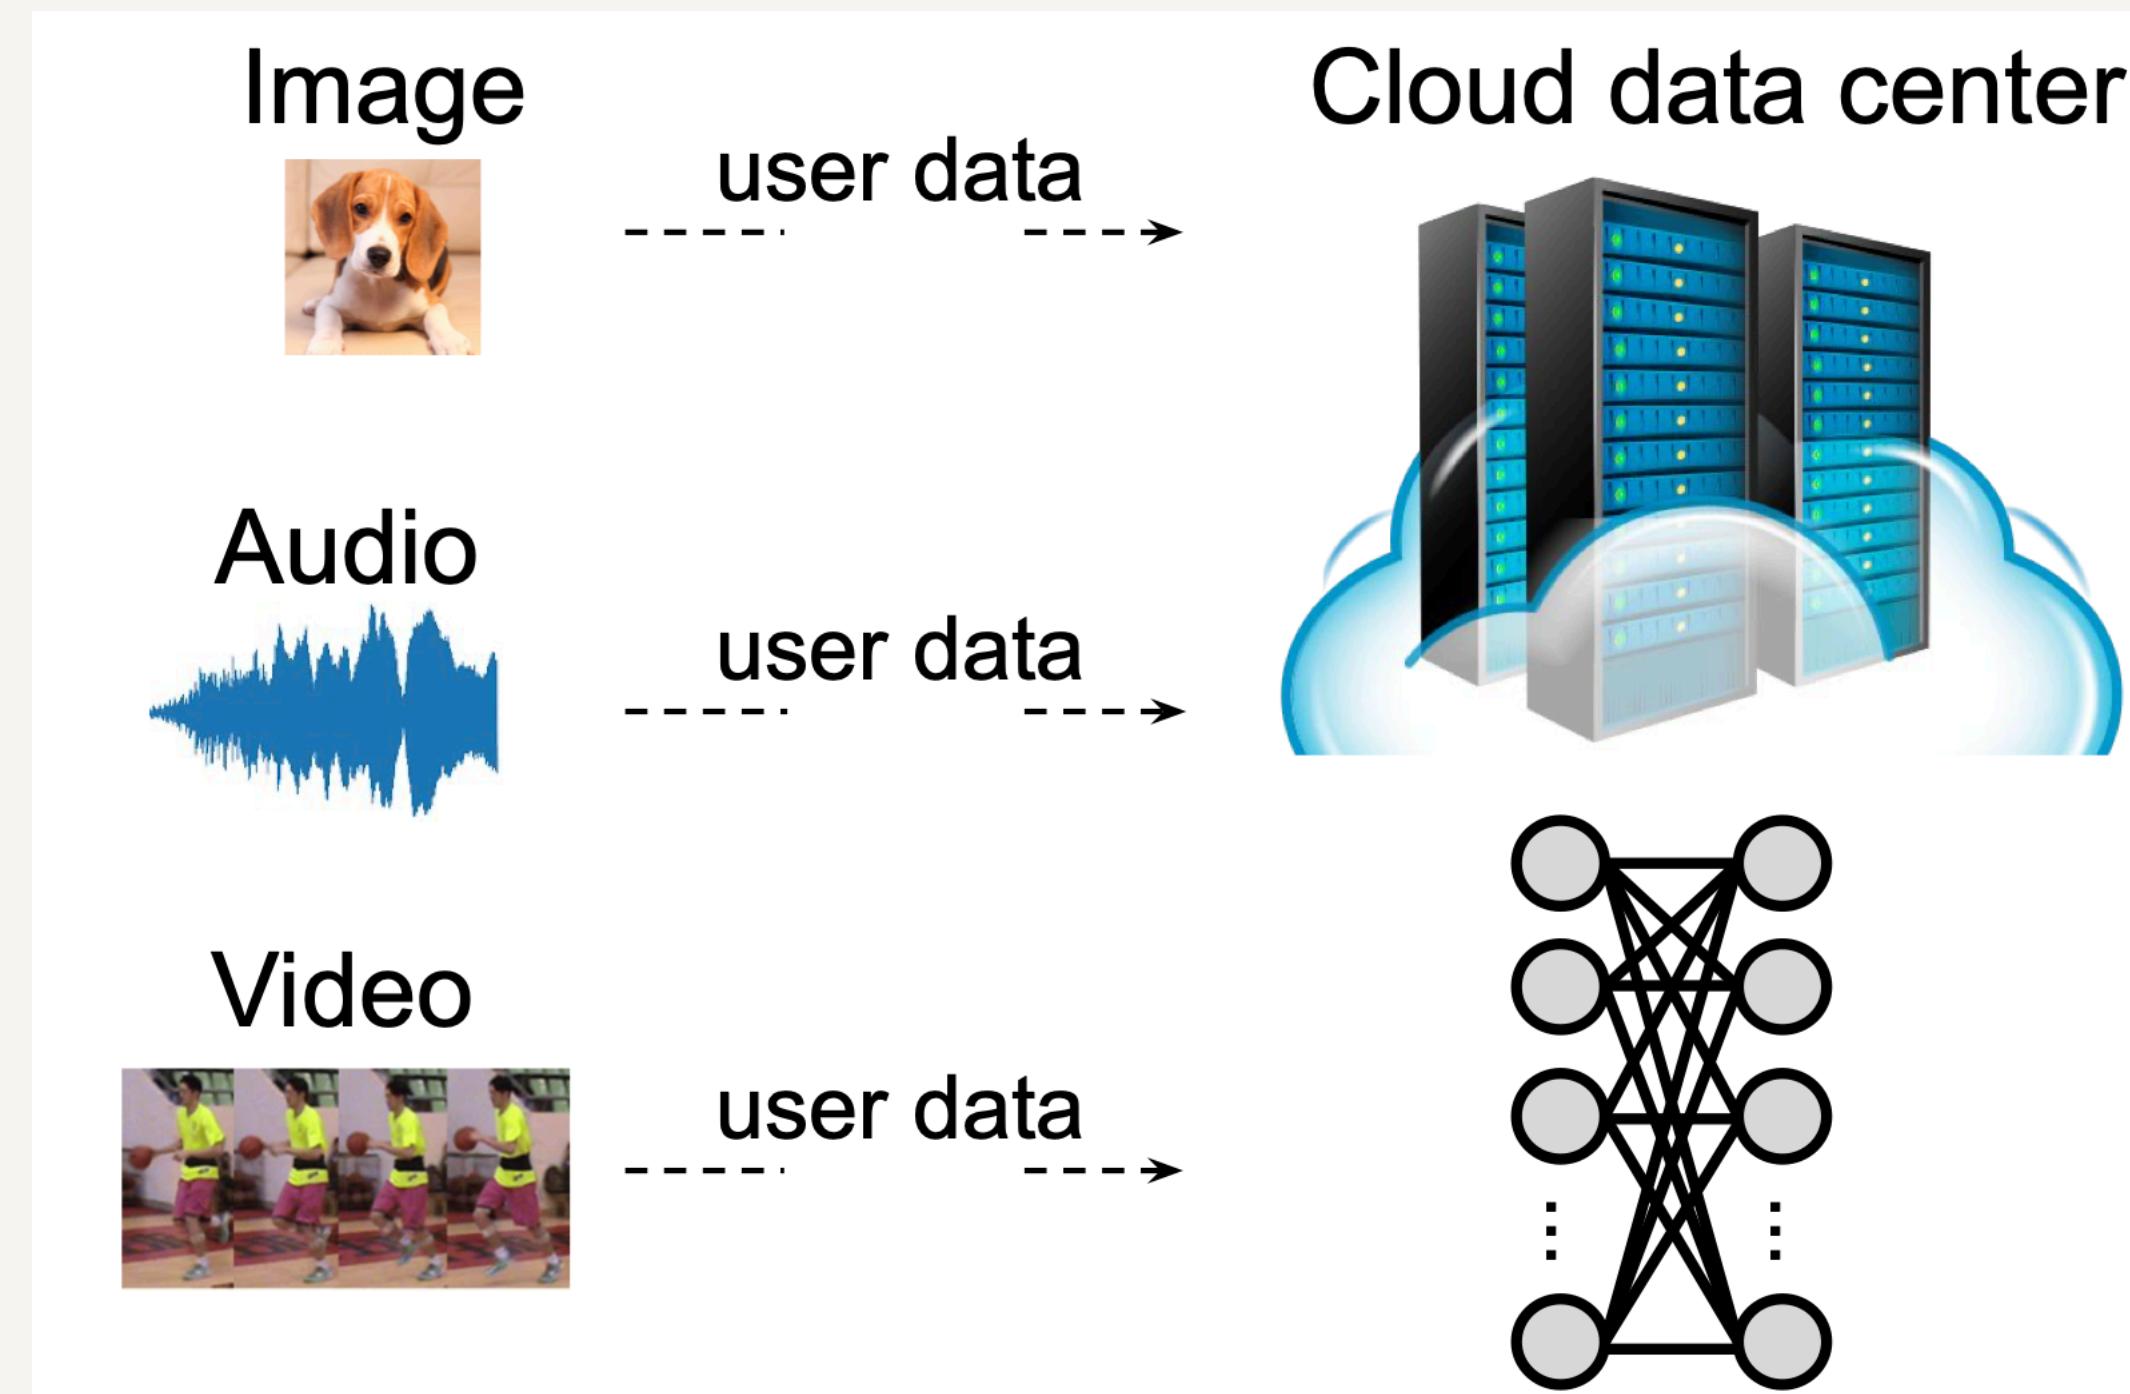


# **Distributed Inference with Deep Learning Models across Heterogeneous Edge Devices**

Chenghao Hu, Baochun Li  
University of Toronto

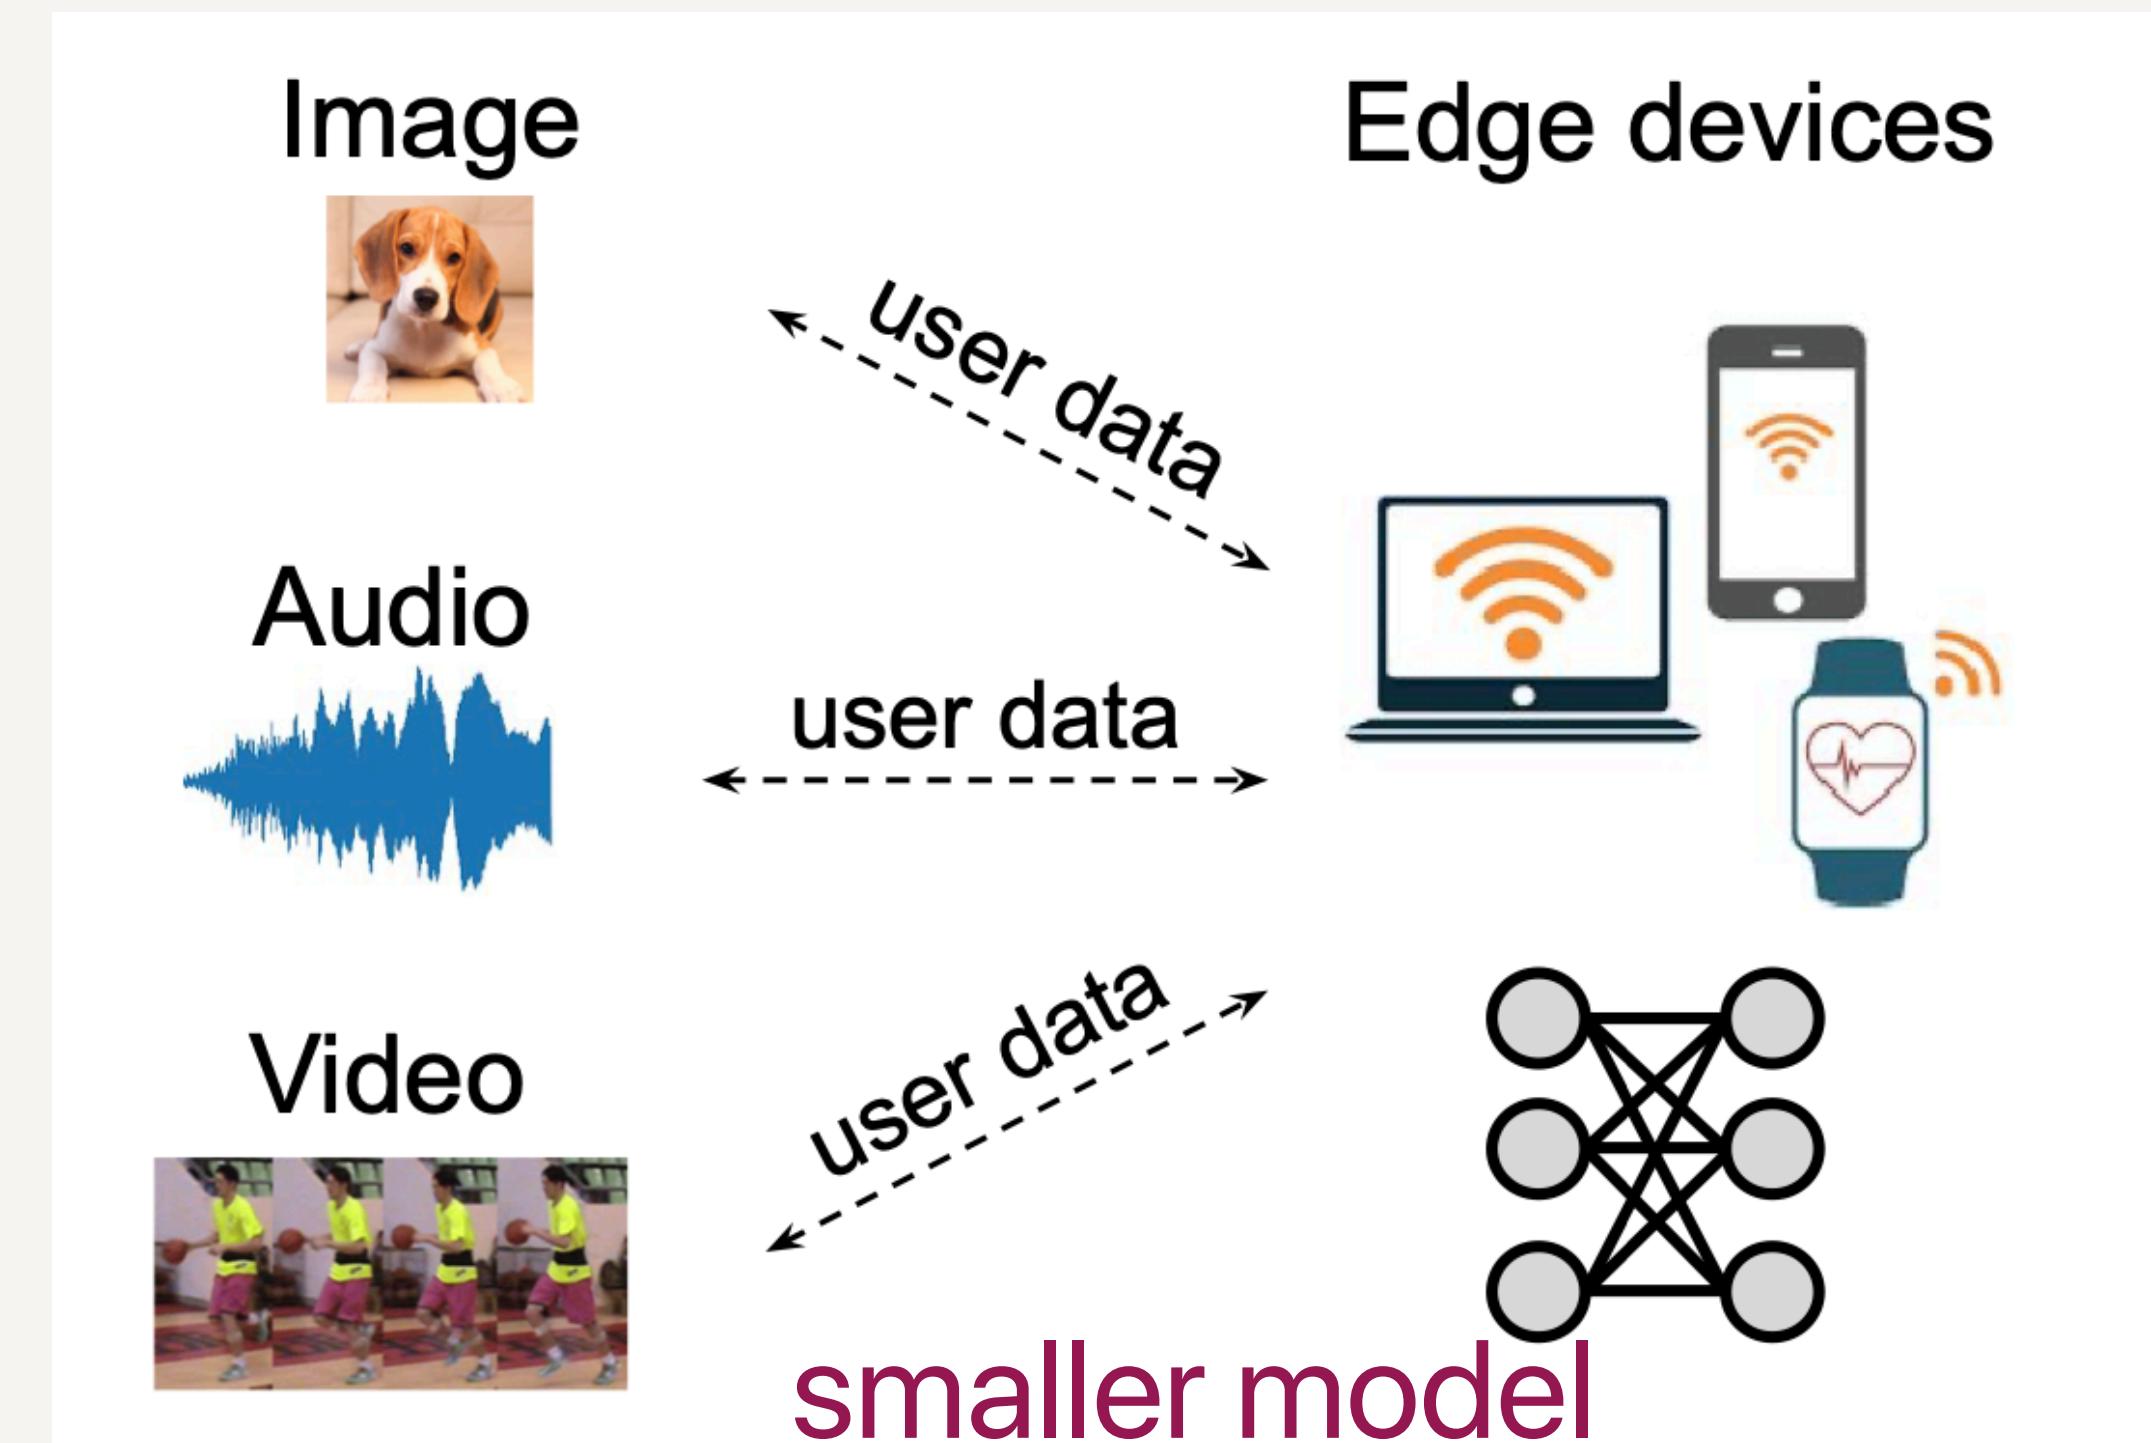
# Model Deployment: Cloud or Edge?

Cloud -> transmission overhead and privacy issue



# Model Deployment: Cloud or Edge?

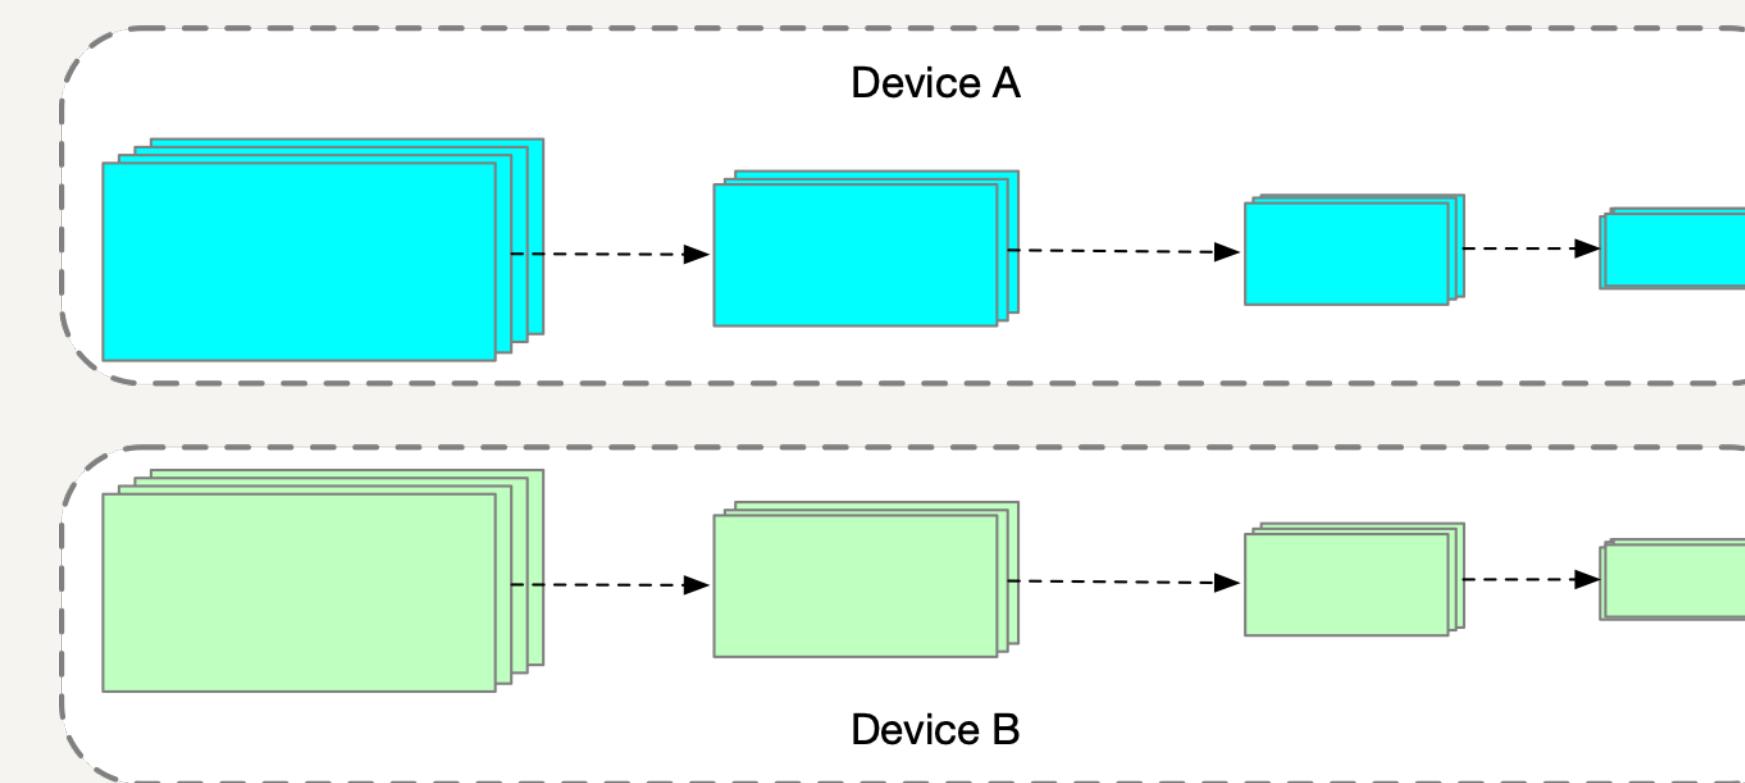
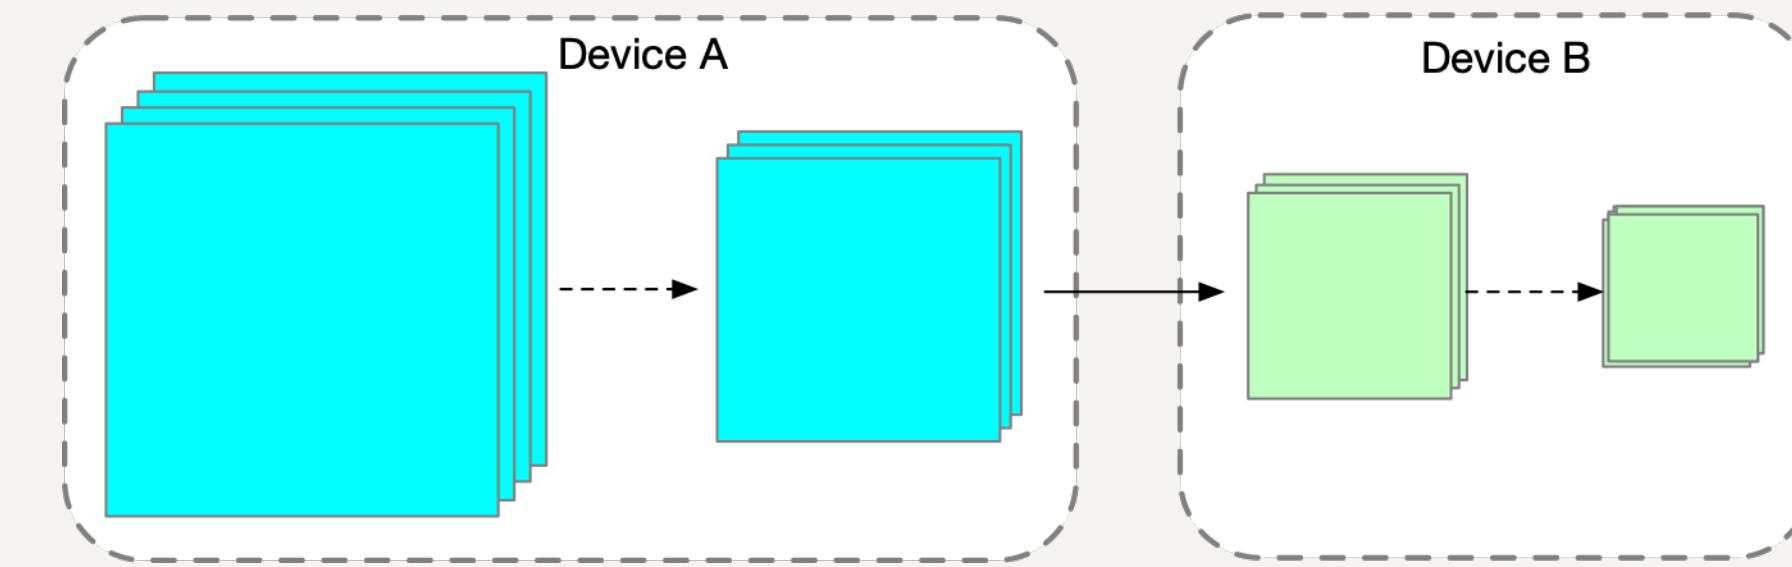
Edge -> Limited computation capacity leads to high latency



Another idea: distribute the inference workload for acceleration

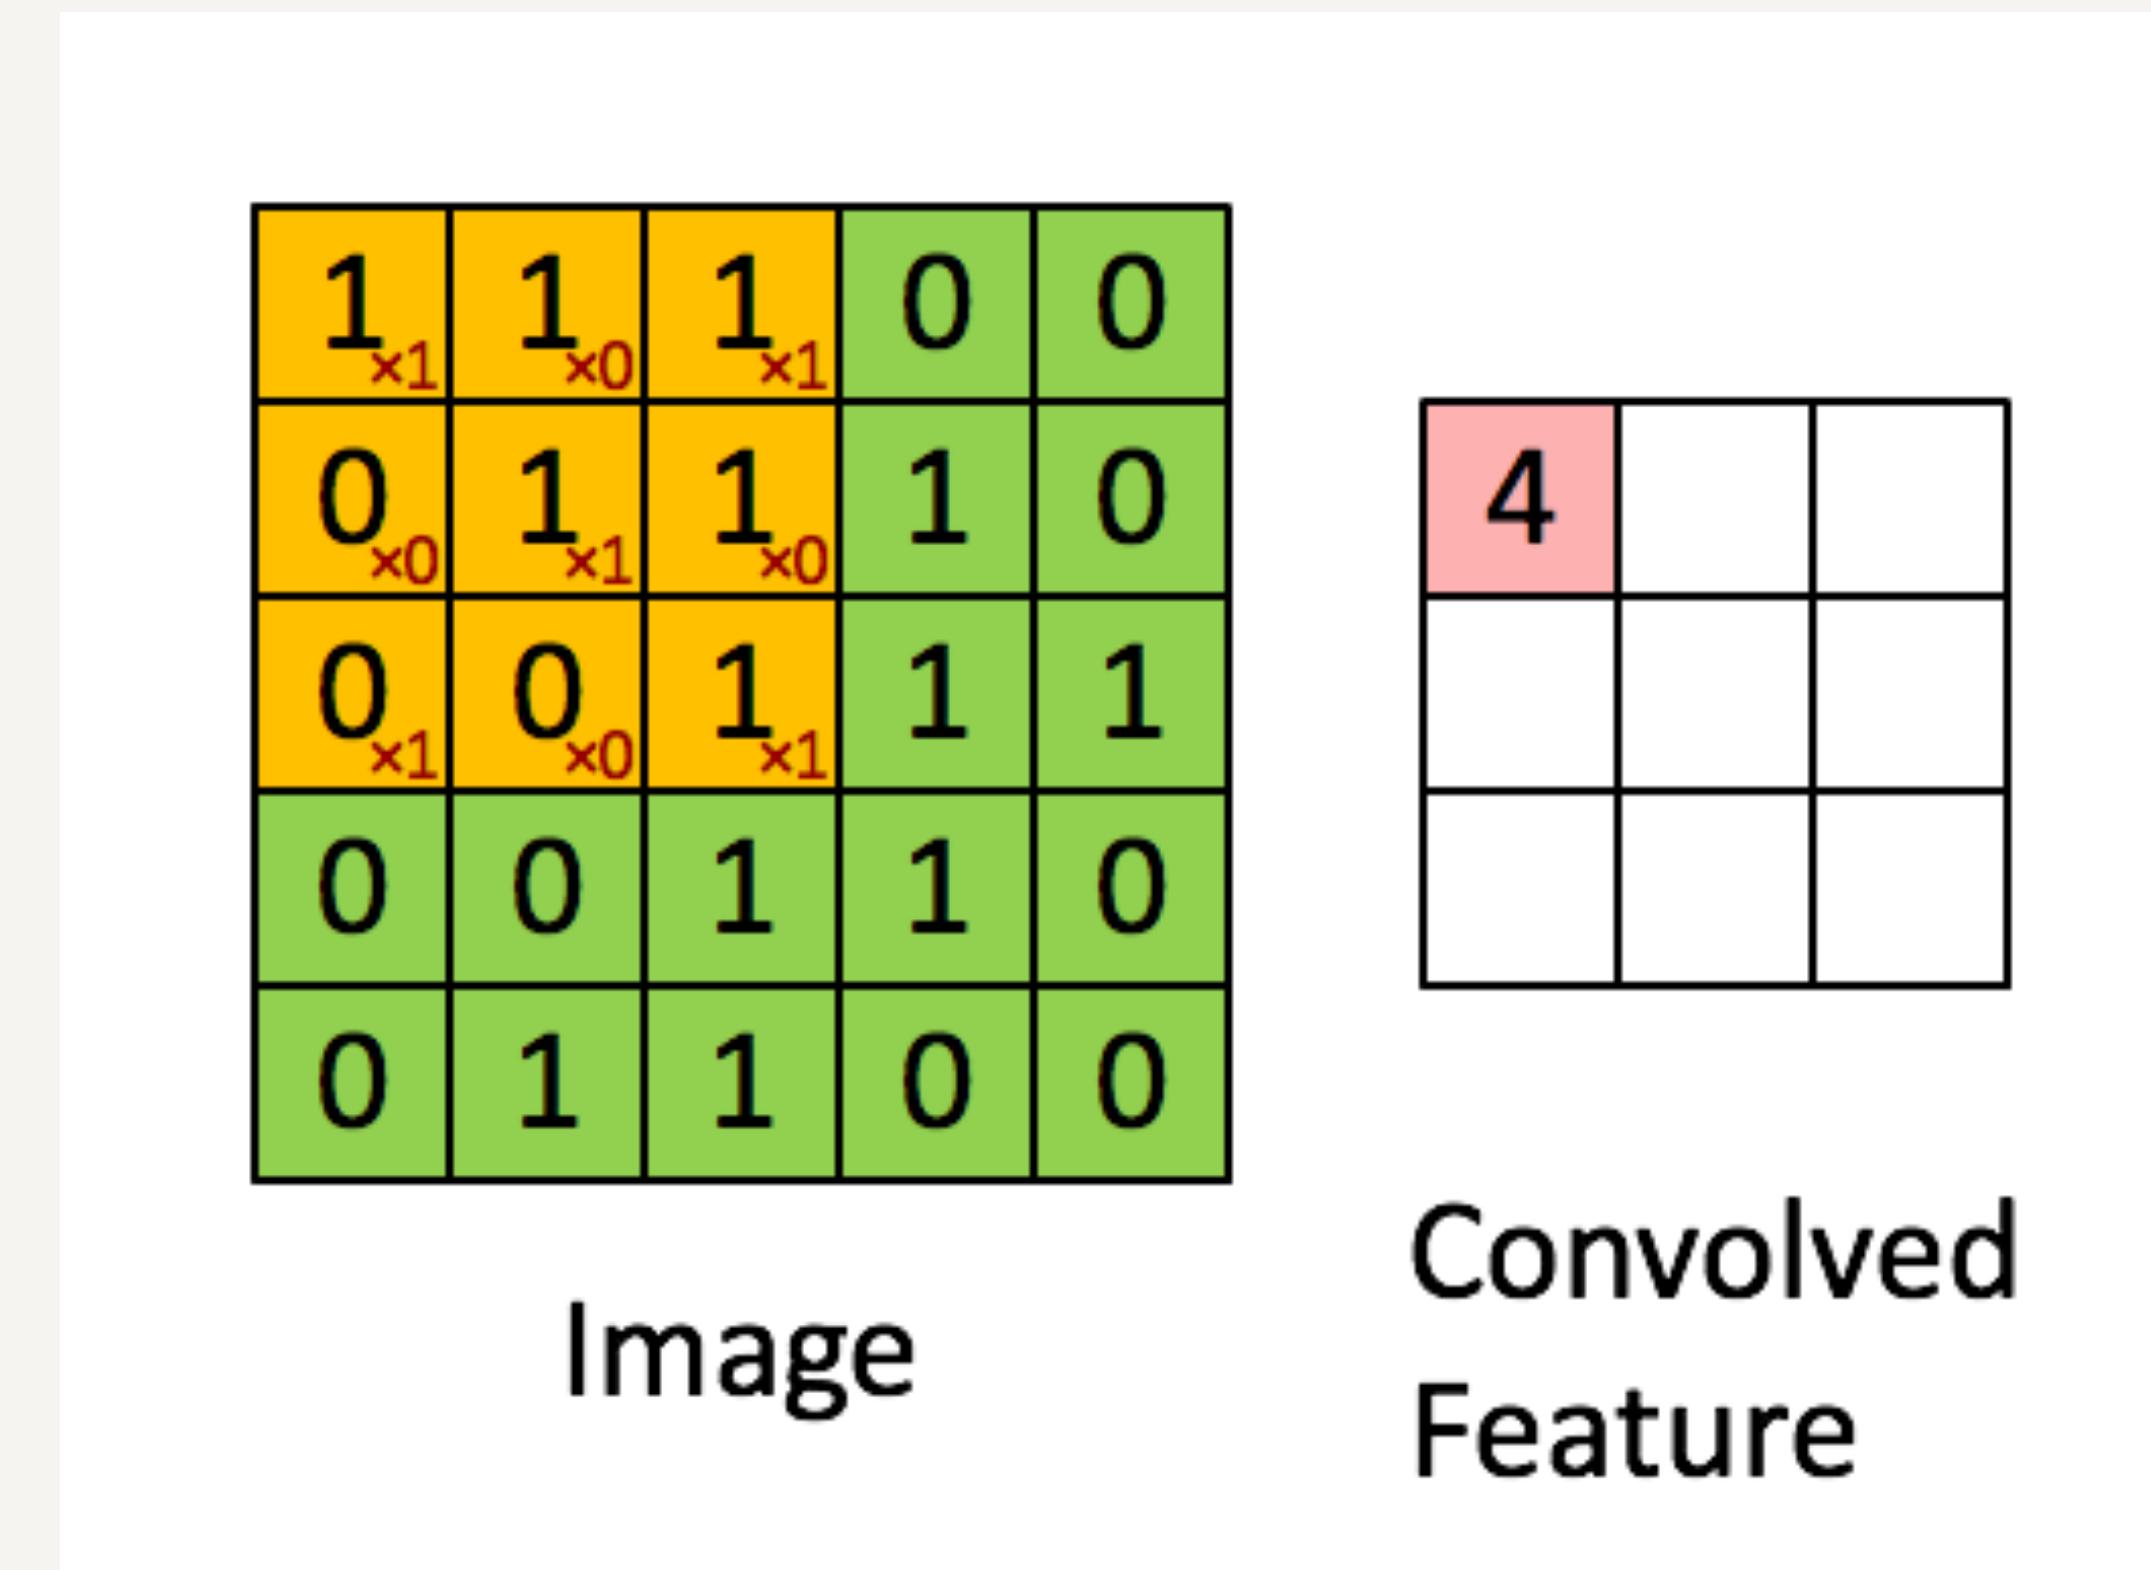
# How to partition the inference workload

- ▶ Sequential partition
  - ▶ Partition the model *layer-wise*
  - ▶ The computation resources are underutilized
- ▶ Parallel partition
  - ▶ Parallel paths executed simultaneously



# Convolution Operation

“Sliding window” applied on the image step by step



# Model Partition - Single Conv Layer

Step 1: decide the range of output partition

Step 2: calculate the range of the required input

Step 3: feed the input partition to the convolution layer

4	3	4

Output Range

1	1	1 <sub>x1</sub>	0 <sub>x0</sub>	0 <sub>x1</sub>
0	1	1 <sub>x0</sub>	1 <sub>x1</sub>	0 <sub>x0</sub>
0	0	1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>
0	0	1	1	0
0	1	1	0	0

Required Input Range

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

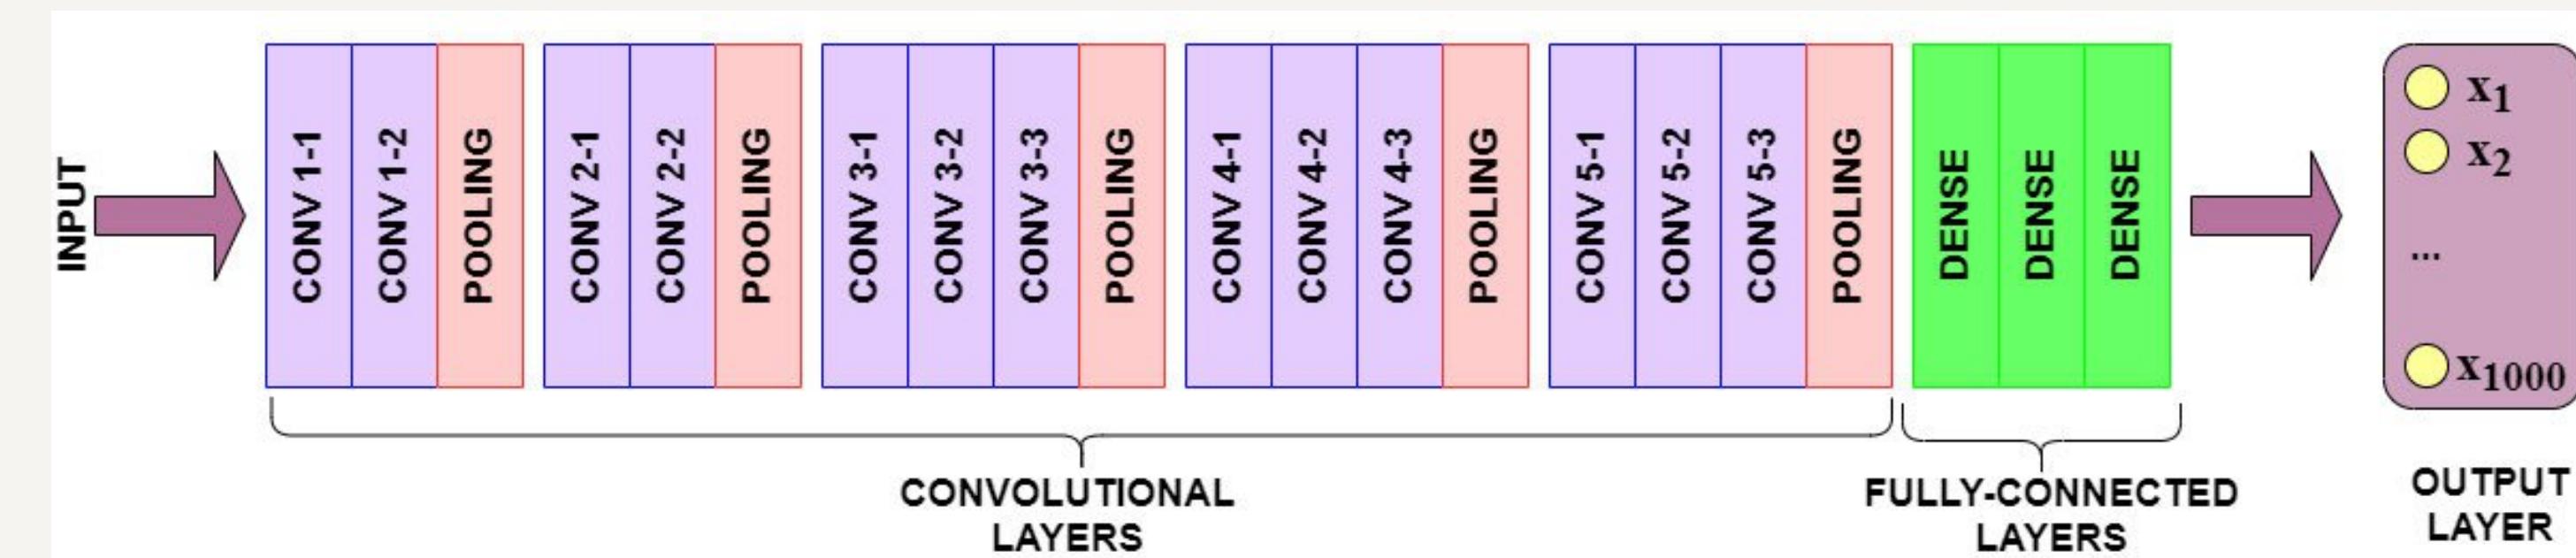
4		

Convolved Feature

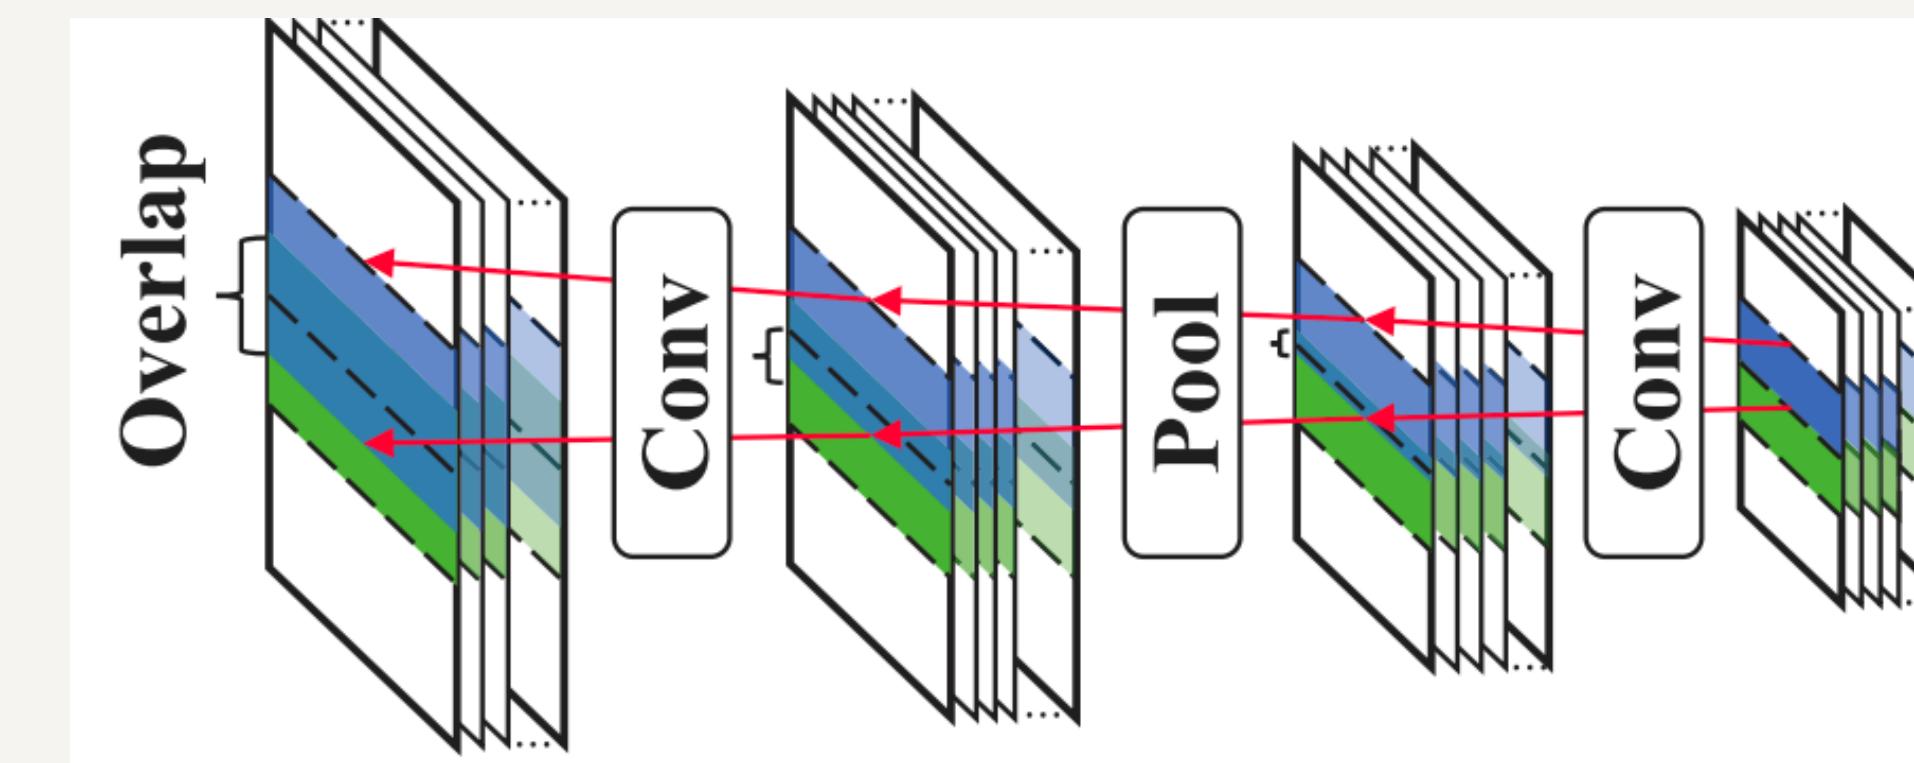
Partitioned Computation

# Model Partition - Chained Layers

Chain structured model, e.g., VGG-16



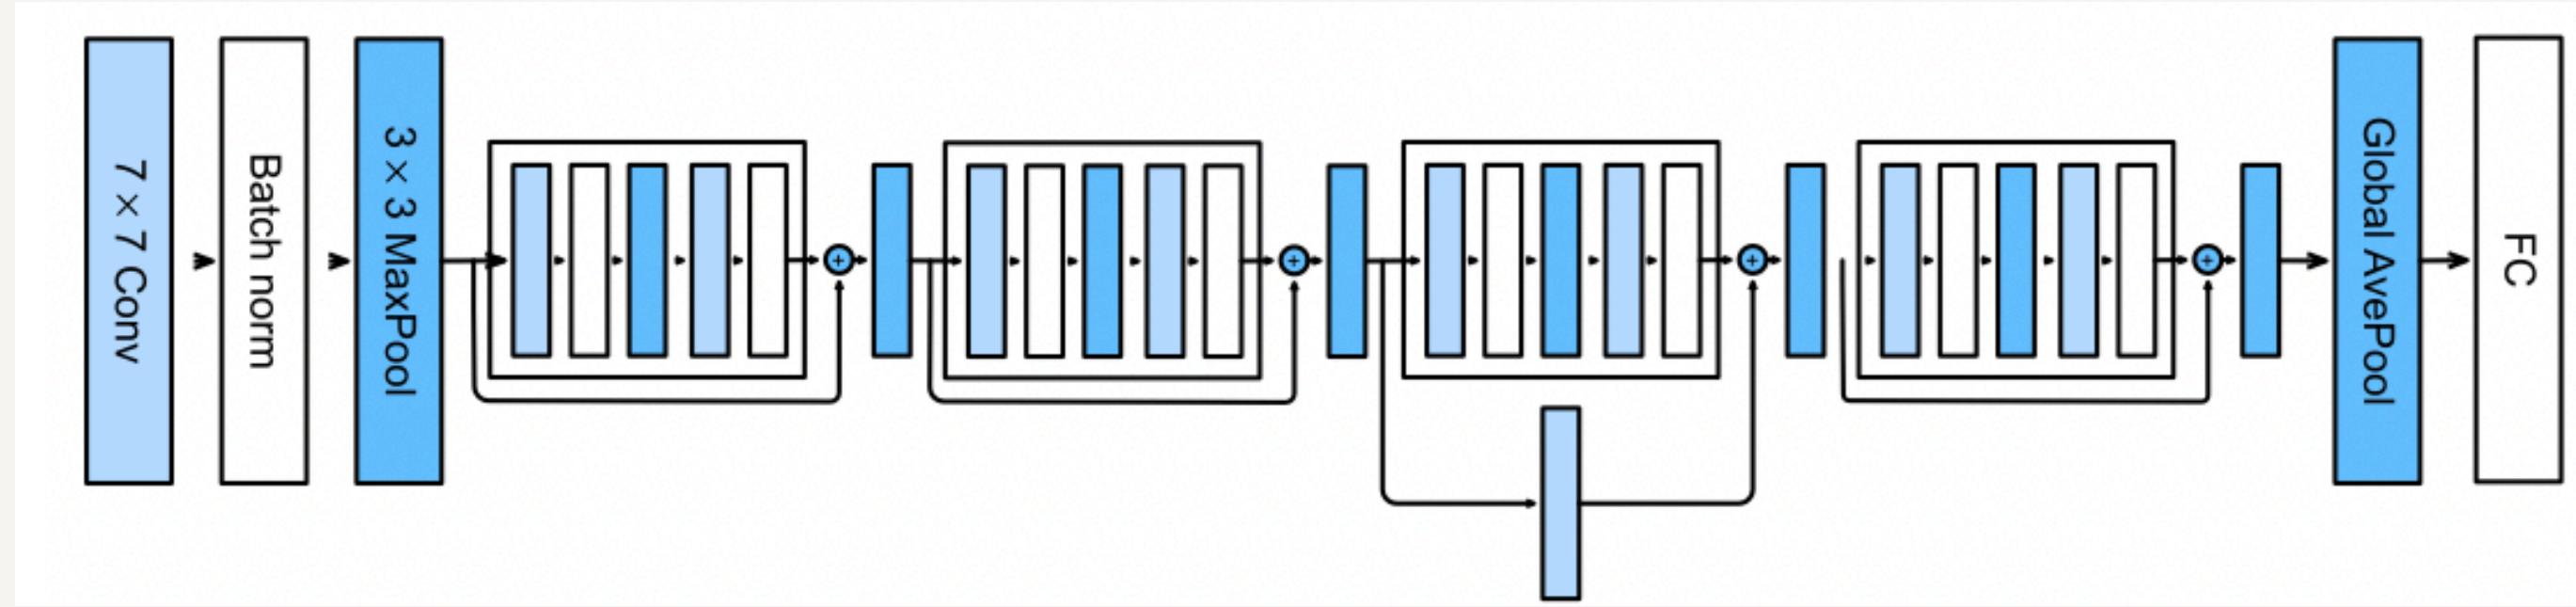
Trace all the way back to the first layer



Existing Solution: DeepThings

# Model Partition - Computation Graph

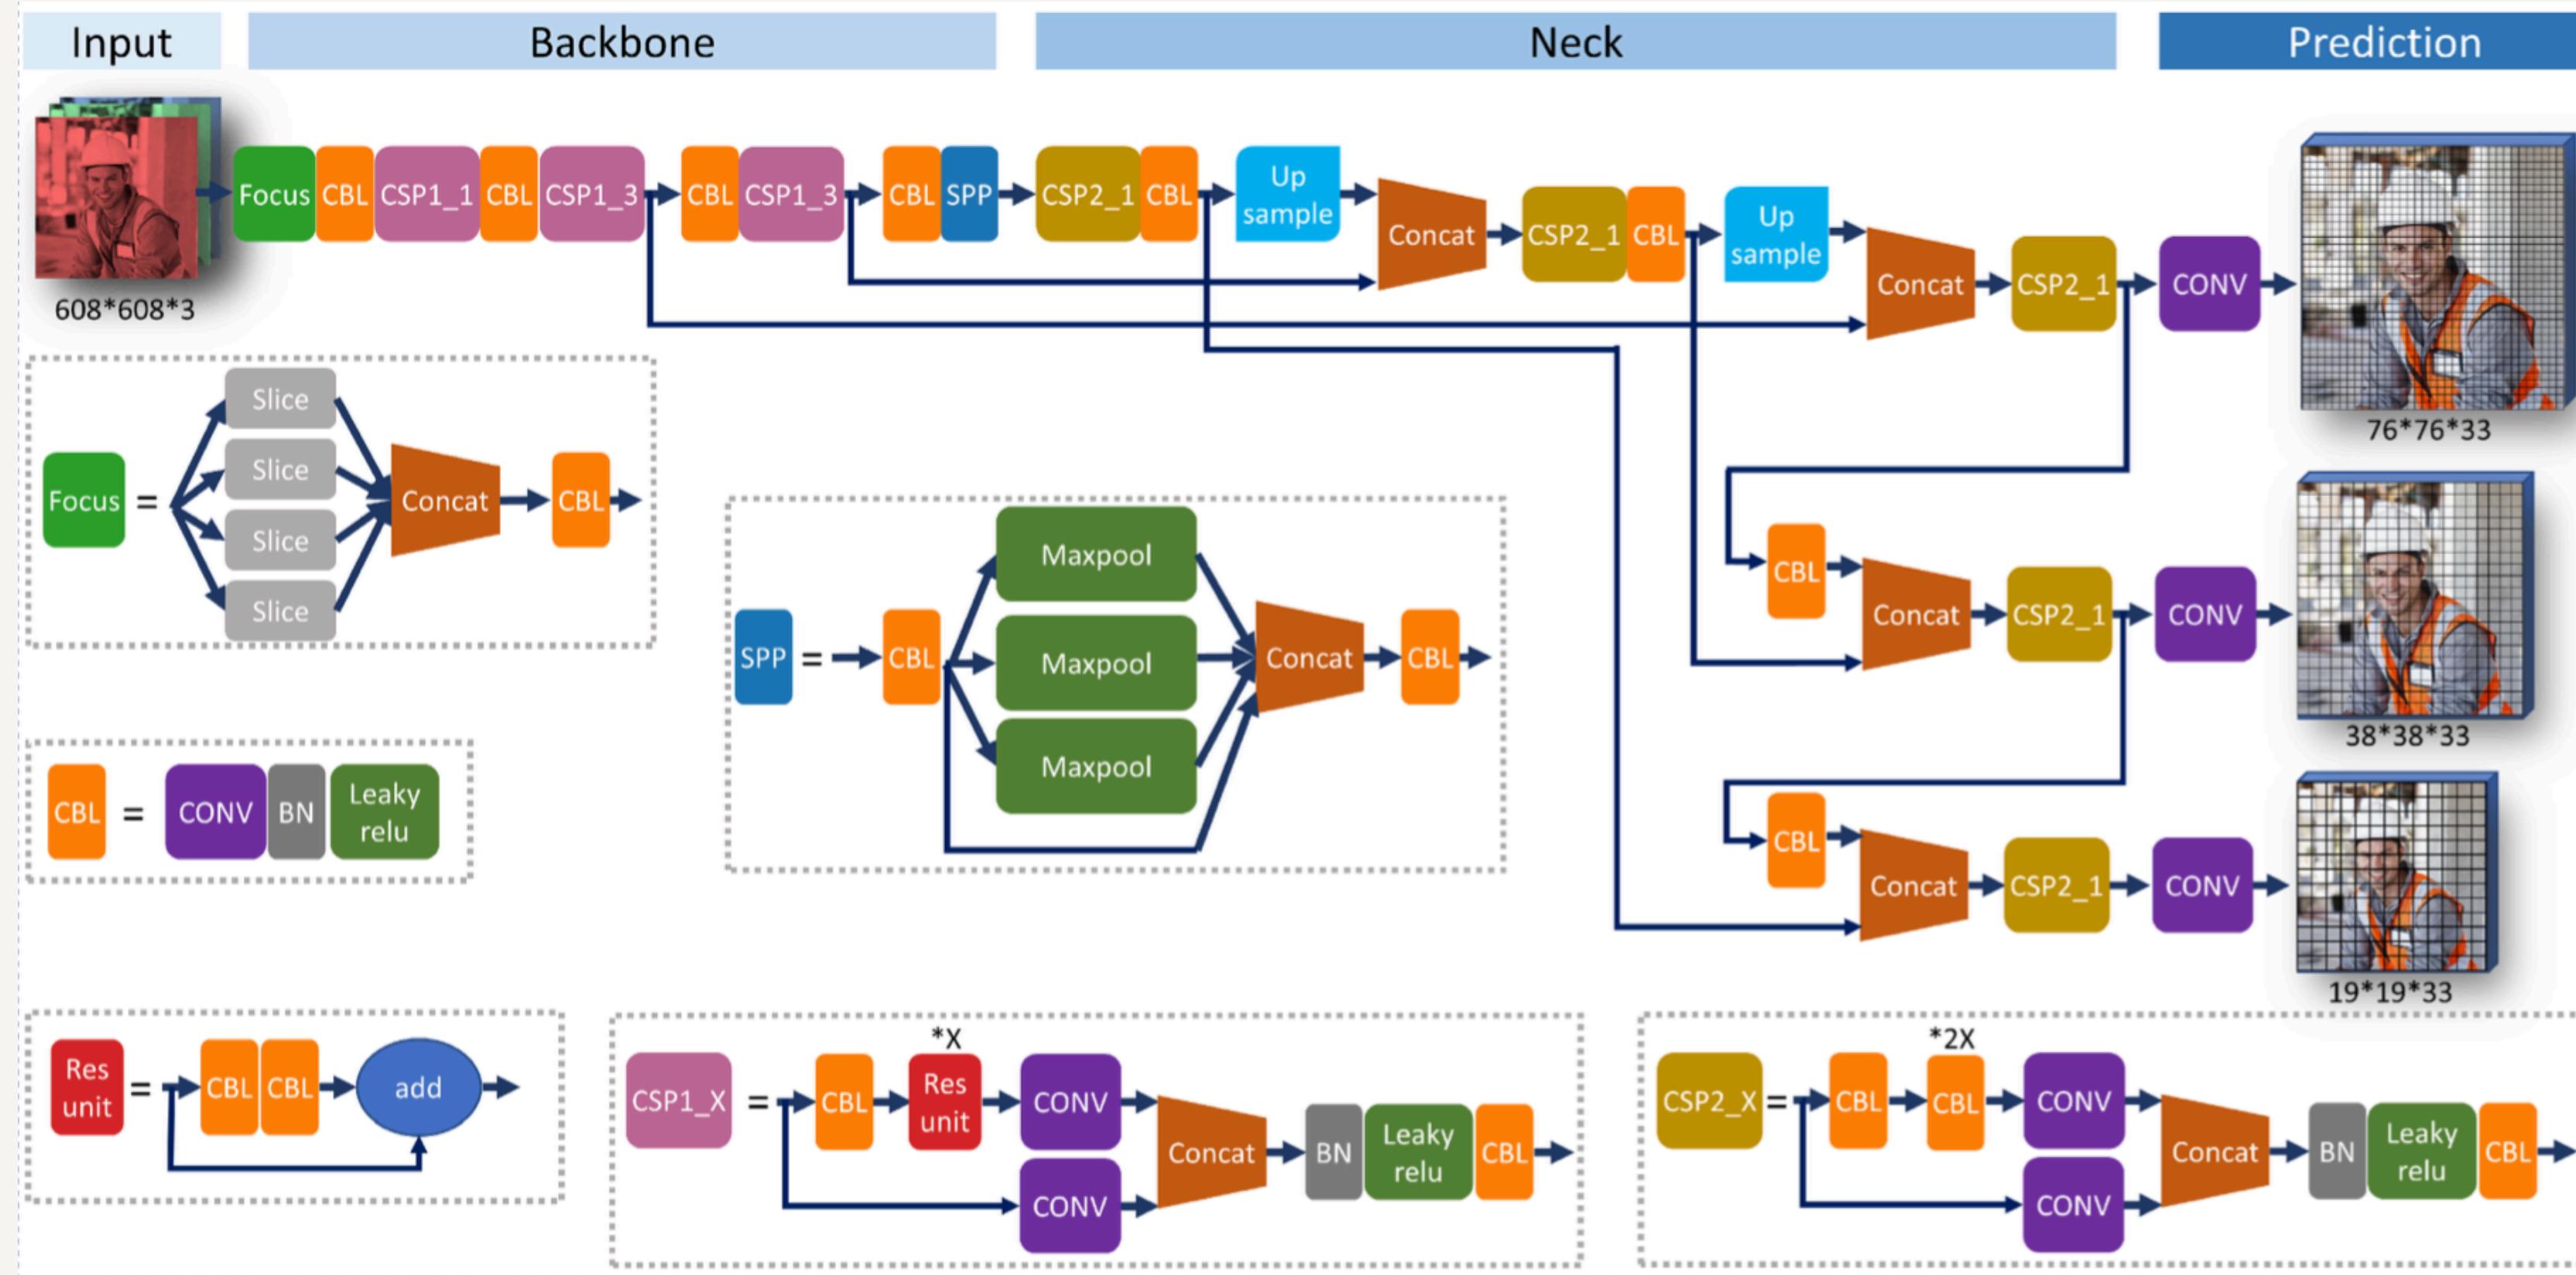
Directed Acyclic Graph (DAG) structured model, e.g., ResNet



Some computation graphs can be easily turned into a chain, and manually fix the layer dependency.

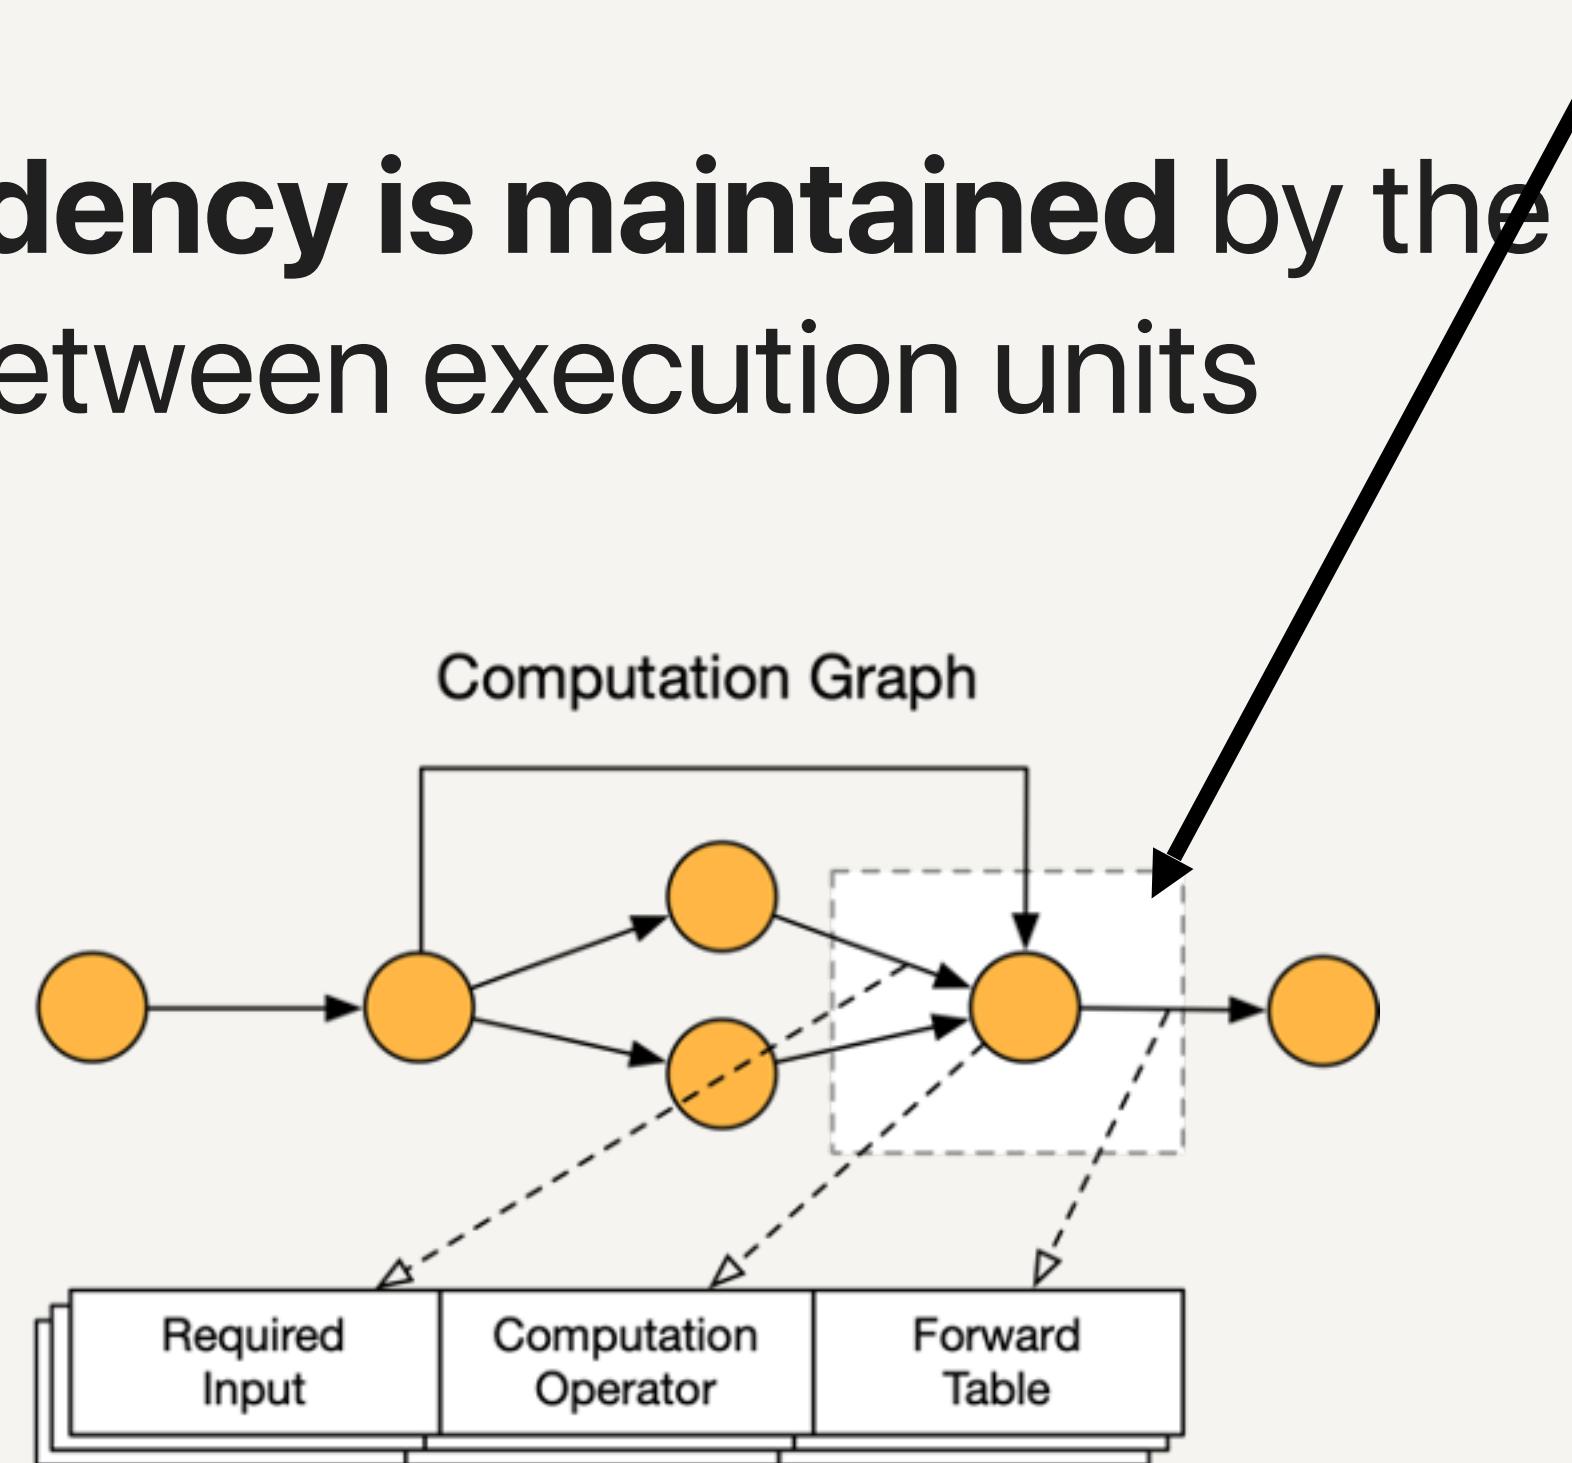
# Model Partition - Computation Graph

Some not... like YoloV5



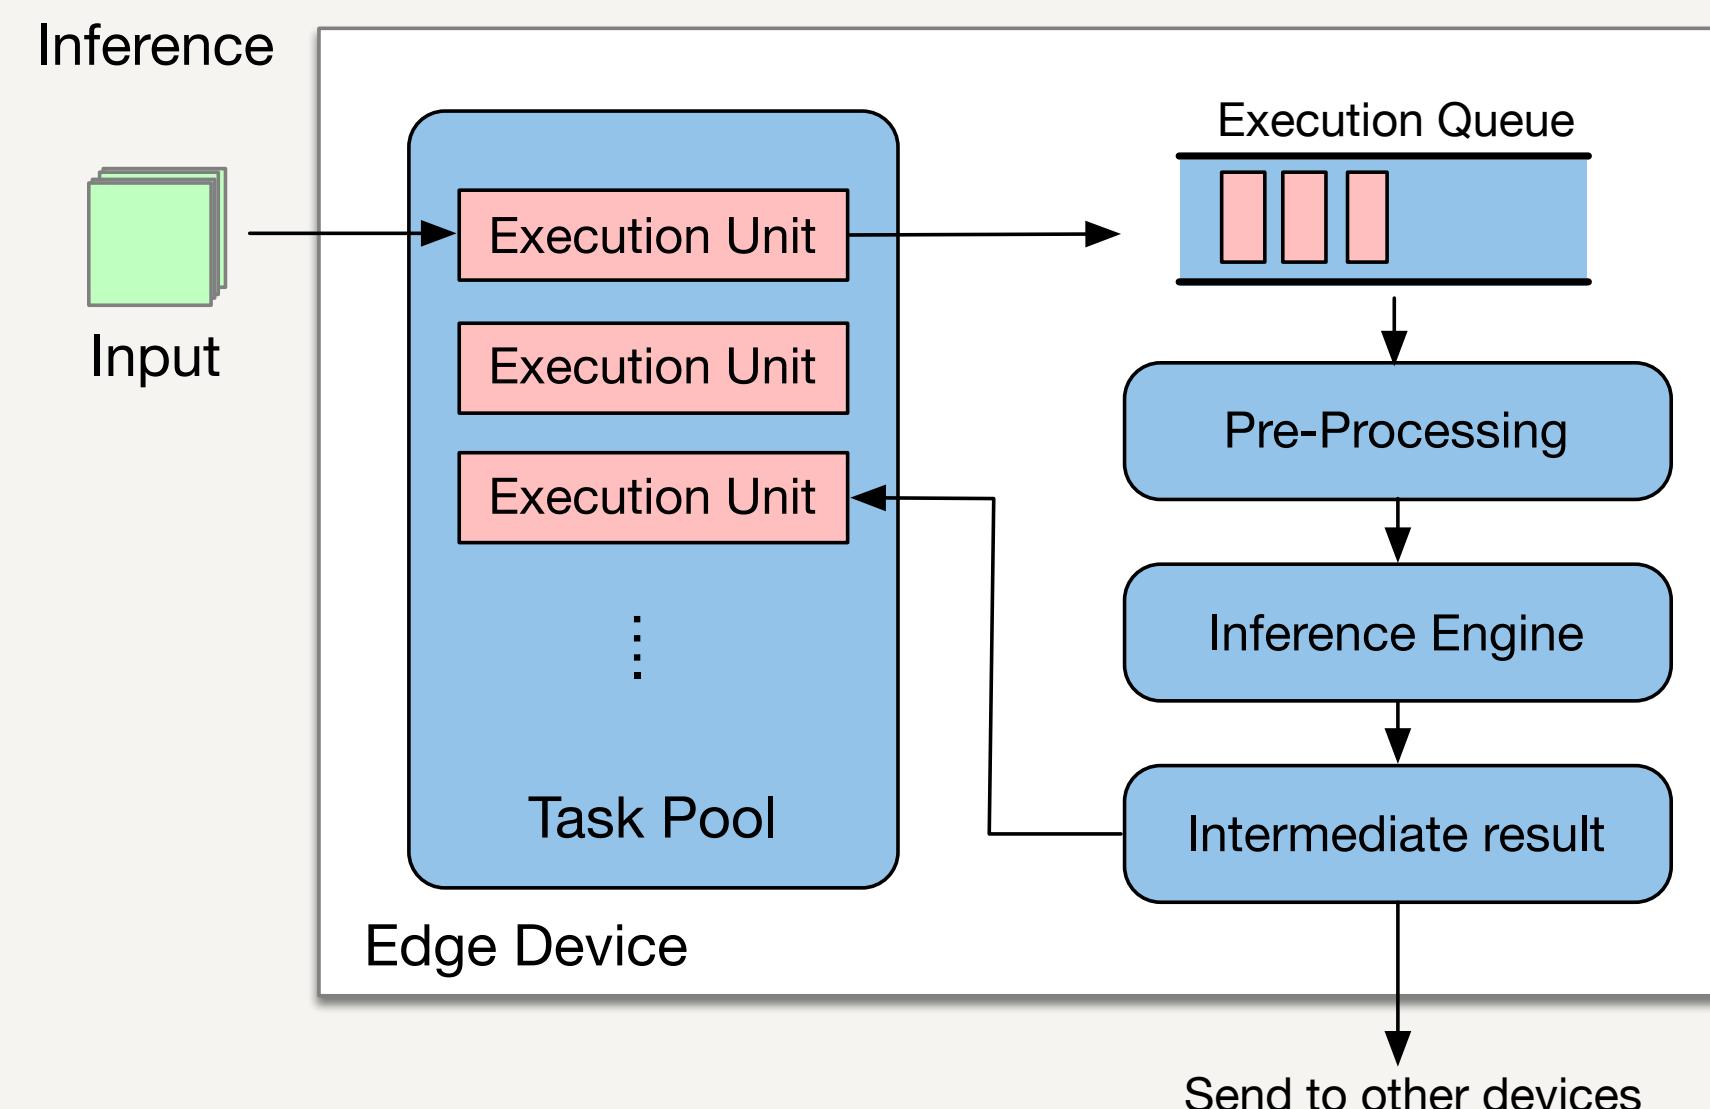
# Our Design - EdgeFlow Overview

- ▶ Setup Phase
  - ▶ partition the computation graph into **execution units**
  - ▶ The **layer dependency is maintained** by the communication between execution units



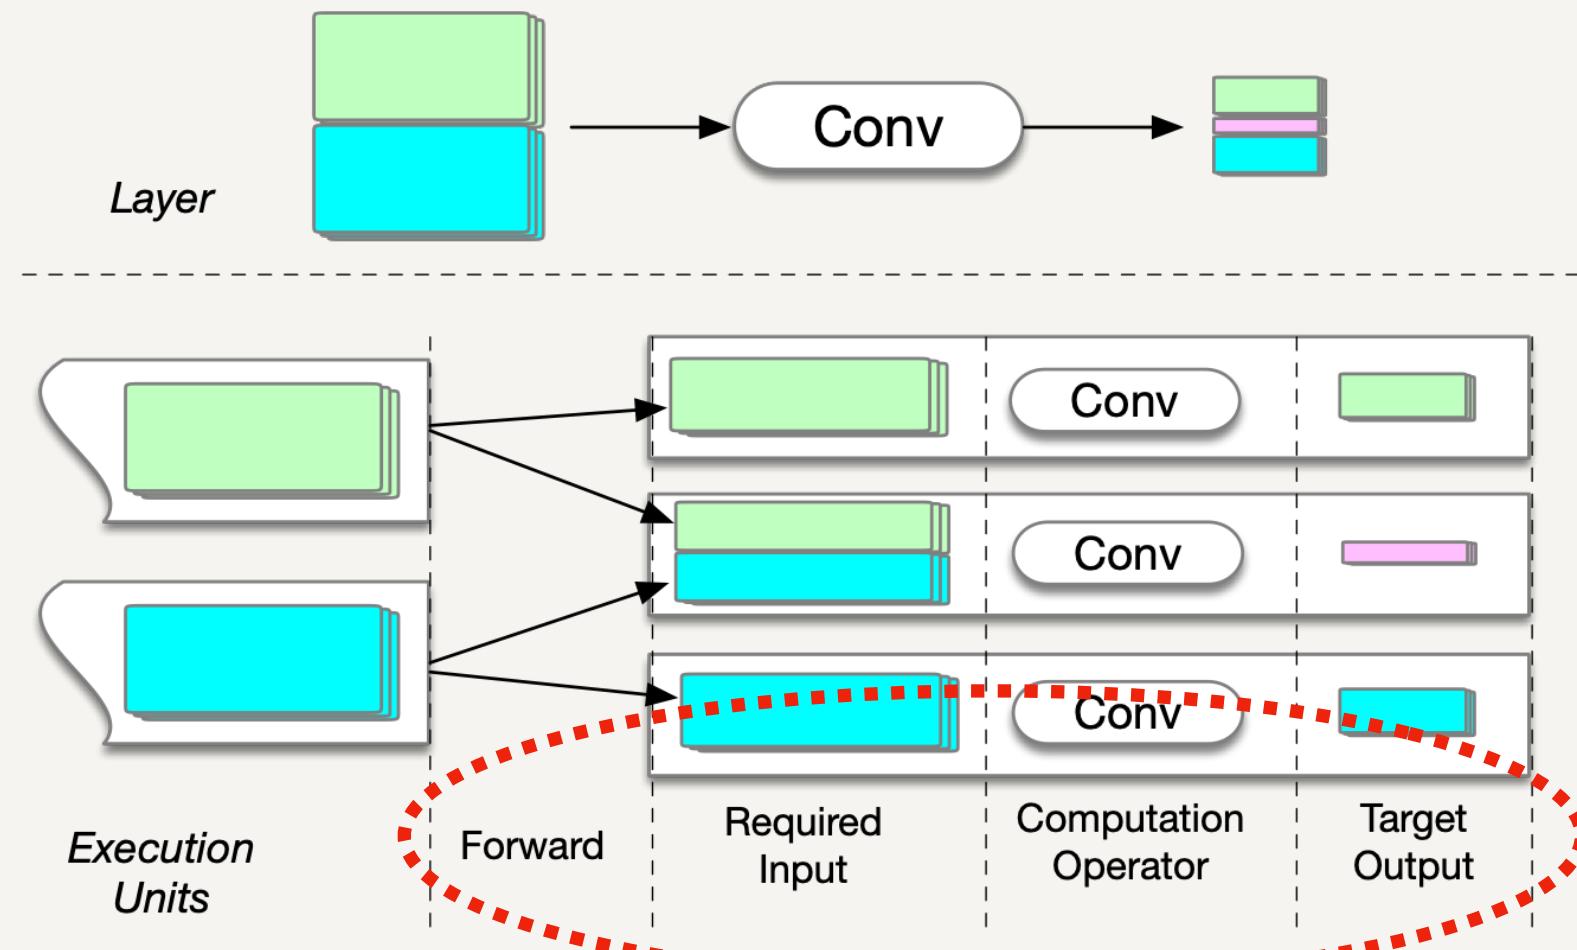
# Our Design - EdgeFlow Overview

- ▶ Inference Phase
  - ▶ Execution units collaboratively finish the inference
  - ▶ *Equivalent* result as computed on a single device



# Model Partitioning

- ▶ Layer partitioning
  - ▶ Each execution unit computes **part** of the output of this layer
  - ▶ Calculate the **required part** of the input needed to complete the computation task
  - ▶ Update the forward table of preceding execution units



Problem: how to find the optimum partition scheme?

# Problem Formulation

- ▶ Assume  $n$  available devices, output features of current layer ranges from row 0 to row  $H$

- ▶ The partition decision variables can be expressed as an integer vector

$$\mathbf{x} = (x_0, x_1, \dots, x_n)$$

- ▶ device  $i$  computes output ranges from row  $x_{i-1} + 1$  to row  $x_i$

$$x_i \in \mathbf{Z}^+, i = 0, 1, \dots, n$$

$$x_0 = 0, x_n = H$$

$$x_0 \leq x_1 \leq \dots \leq x_n$$

# Problem Formulation

- ▶ Objective: finish time of the current layer  $l$ 
  - ▶  $T_{l,i}$  denote the time that device  $i$  finish its partition of layer  $l$
- ▶ The optimization problem can expressed as

$$\begin{aligned} \min_x \quad & \max(T_{l,1}, T_{l,2}, \dots, T_{l,n}) \\ s.t. \quad & x_i \in \mathbf{Z}^+, i = 0, 1, \dots, n \\ & x_0 = 0, x_n = H \\ & x_0 \leq x_1 \leq \dots \leq x_n \end{aligned}$$

# Problem Formulation

- $T_{l,i}$  estimation: transmission time + computation time

$$T_{l,i} = t_{\text{trans}}(i; l) + t_{\text{comp}}(i; l)$$

- Computation time can be estimated with a pre-trained linear regression model

$$t_{\text{comp}}(i) = Y_i(x_i - x_{i-1}; l).$$

↑  
number of rows to compute  
↓  
layer settings

- Transmission time can be estimated by

$$t_{\text{trans}}(i; l) = \max_{j \in \{1, \dots, n\}, m \in M} 1_{\{p_{m,i,j} > 0\}} (T_{m,j} + \frac{p_{m,i,j}}{B_{i,j}})$$

The layers required by current layer

valid when the transmission is  
positive

↑  
finish time of layer m at device j  
↑  
wait all the transmission done  
↓  
Transmission from j to i

# Problem Transformation

- ▶ Original problem

$$\min_x \max(T_{l,1}, T_{l,2}, \dots, T_{l,n})$$

$$s.t. \quad x_i \in \mathbf{Z}^+, i = 0, 1, \dots, n$$

$$x_0 = 0, x_n = H$$

$$x_0 \leq x_1 \leq \dots \leq x_n$$

- ▶ Step 1: introducing auxiliary variable and relax the integer constraint

$$\min_{x, \lambda} \lambda$$

$$s.t. \quad T_{l,i} \leq \lambda, i \in \{1, \dots, n\}$$

$$x_i \in \mathbf{Z}^+, i = 0, 1, \dots, n$$

$$x_0 = 0, x_n = H$$

$$x_0 \leq x_1 \leq \dots \leq x_n$$

# Problem Transformation

- ▶ Step 2: removing the indicator function

$$1_{\{p_{m,i,j} > 0\}}(T_{m,j} + \frac{p_{m,i,j}}{B_{i,j}}) + Y_i(x_i - x_{i-1}; l) \leq \lambda$$

The finish time of layer  $m$  at different device  
should be roughly the same

As long as  $T_{m,j}$  is not greater than other  
required devices, the constraint is loose, and  
won't affect the result.

# Problem Transformation

- ▶ Step 3: re-express the transmission size

$p_{m,i,j}$  is the overlapping area between the **required input range**  $(s_i, e_i)$

and the **output range**  $(x_{m,j-1}, x_{m,j})$  of layer  $m$  at device  $j$

$$\begin{aligned} p_{m,i,j} &= \min(e_i, x_{m,j}) - \max(s_i, x_{m,j-1}) \\ &= \min(e_i - s_i, e_i - x_{m,j-1}, x_{m,j} - s_i, x_{m,j} - x_{m,j-1}) \\ &\quad \downarrow \\ \min_{p_{m,i,j}} \quad &- p_{m,i,j} \end{aligned}$$

$$\begin{aligned} s.t. \quad p_{m,i,j} &\leq e_i - s_i, \quad p_{m,i,j} \leq e_i - x_{m,j-1}, \\ p_{m,i,j} &\leq x_{m,j} - s_i, \quad p_{m,i,j} \leq x_{m,j} - x_{m,j-1}. \end{aligned}$$

# Problem Transformation

- ▶ Linear Programming Approximation

$$\min_{x, \lambda, p} \quad \lambda - \sum_{m,i,j} p_{m,i,j}$$

$$s.t. \quad x_0 = 0, x_n = H$$

$$x_0 \leq x_1 \leq \dots \leq x_n$$

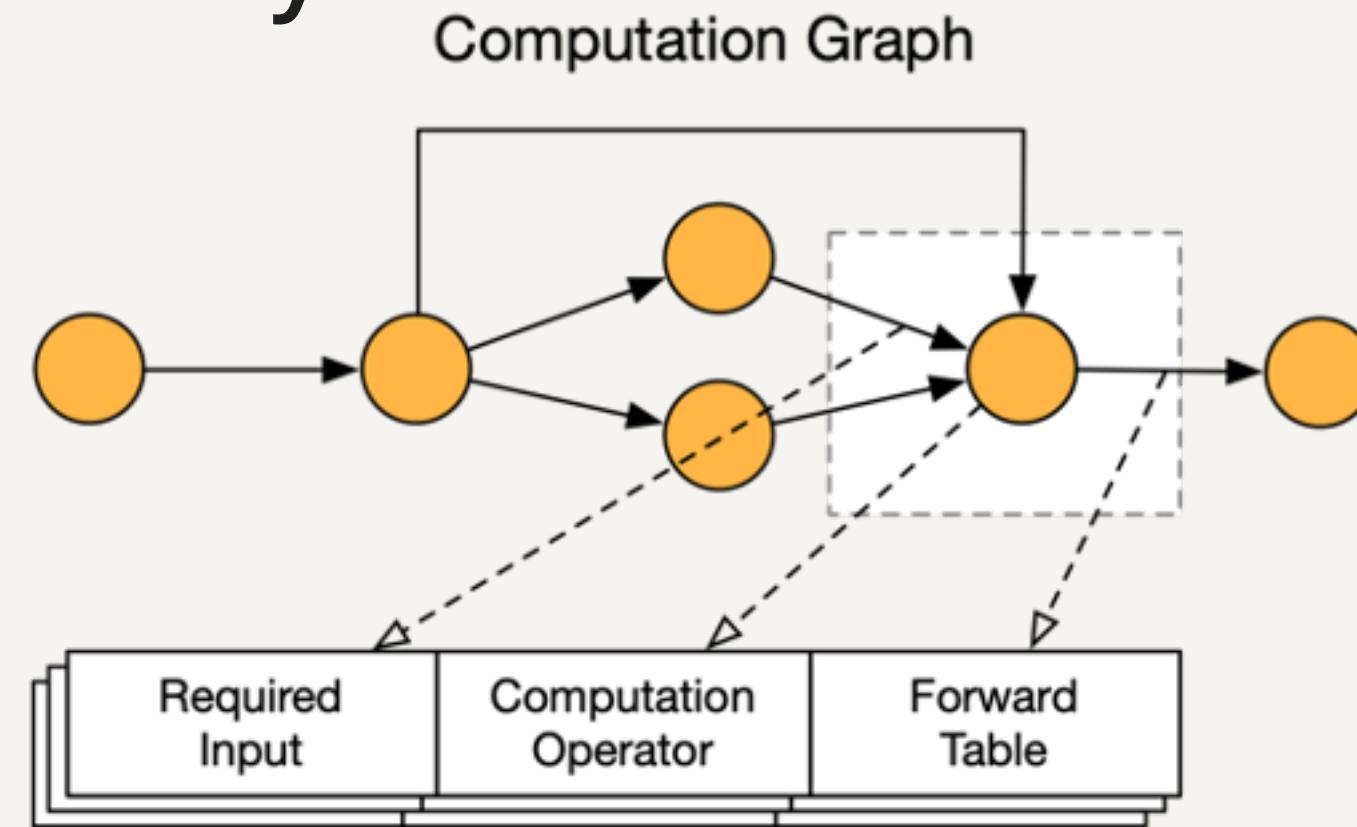
$$T_{m,j} + \frac{p_{m,i,j}}{B_{i,j}} + Y_i(x_i - x_{i-1}; l) \leq \lambda$$

$$p_{m,i,j} \leq e_i - s_i, \quad p_{m,i,j} \leq e_i - x_{m,j-1},$$

$$p_{m,i,j} \leq x_{m,j} - s_i, \quad p_{m,i,j} \leq x_{m,j} - x_{m,j-1}.$$

# Model Partition

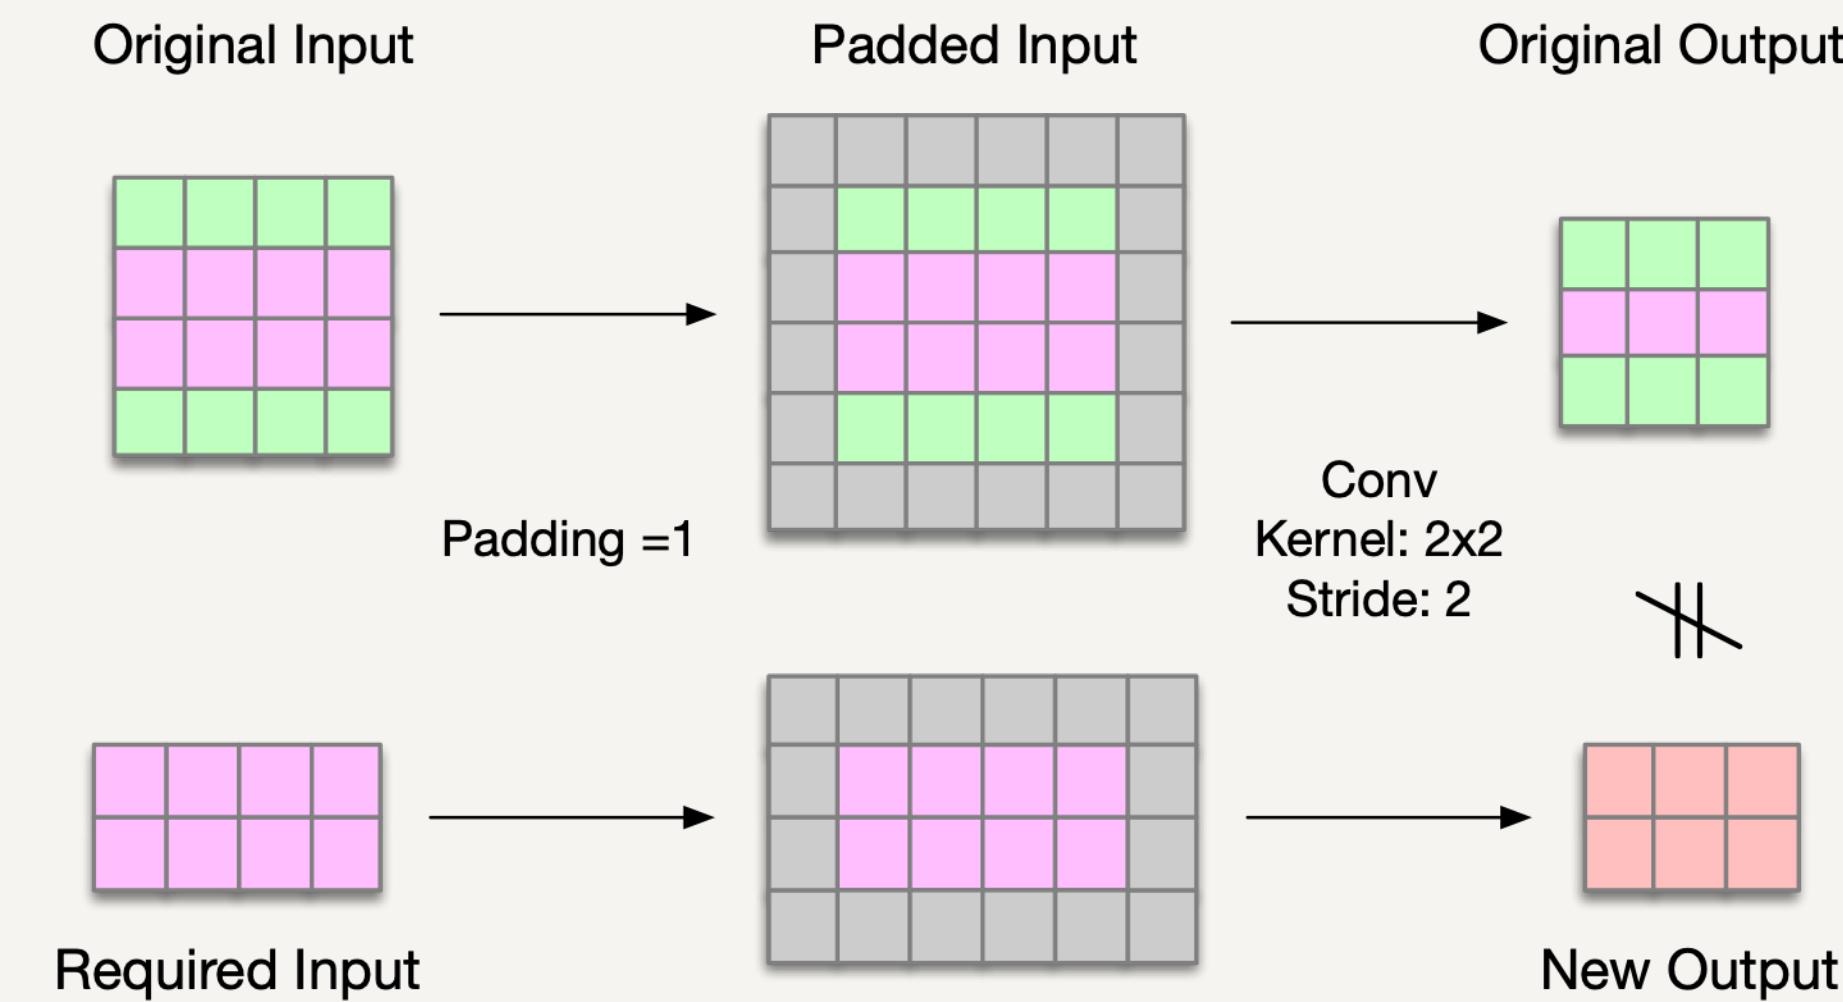
- ▶ Partition the model layer by layer in topological order
- ▶ Solve the LP problem for each layer to obtain the partition scheme
- ▶ The finish time estimation of previous layer becomes a parameter of the optimization problem of the following layers



$$T_{m,j} + \frac{p_{m,i,j}}{B_{i,j}} + Y_i(x_i - x_{i-1}; l) \leq \lambda$$

# Padding Issue

- Directly feeding the input partition to the conv/pool layer may not yield the correct output



# Padding Issue

- ▶ Solution: pre-padding mechanism
  - ▶ Set the padding parameter of conv/pool layer to 0
  - ▶ Manually add paddings when necessary

$$i_s = o_s \times \text{stride} - \text{padding},$$

$$i_e = (o_e - 1) \times \text{stride} + \text{kernel\_size} - \text{padding},$$

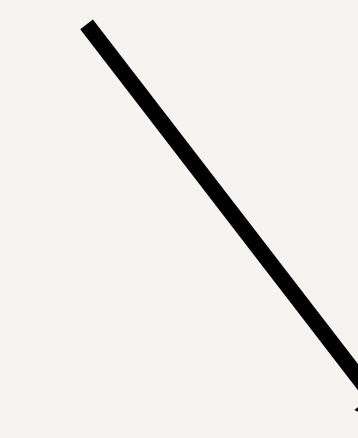
$$\text{upper\_padding} = \begin{cases} -i_s, & i_s < 0 \\ 0, & \text{otherwise} \end{cases},$$

$$\text{bottom\_padding} = \begin{cases} i_e - H_i, & i_e > H_i \\ 0, & \text{otherwise} \end{cases}.$$

# Inference Phase

- ▶ The units will be executed when the input requirements are satisfied
- ▶ The output will be forwarded to fulfill the requirement of next execution unit
- ▶ Intermediate results ***flow*** through execution units to finish the inference

System name : ***EdgeFlow***

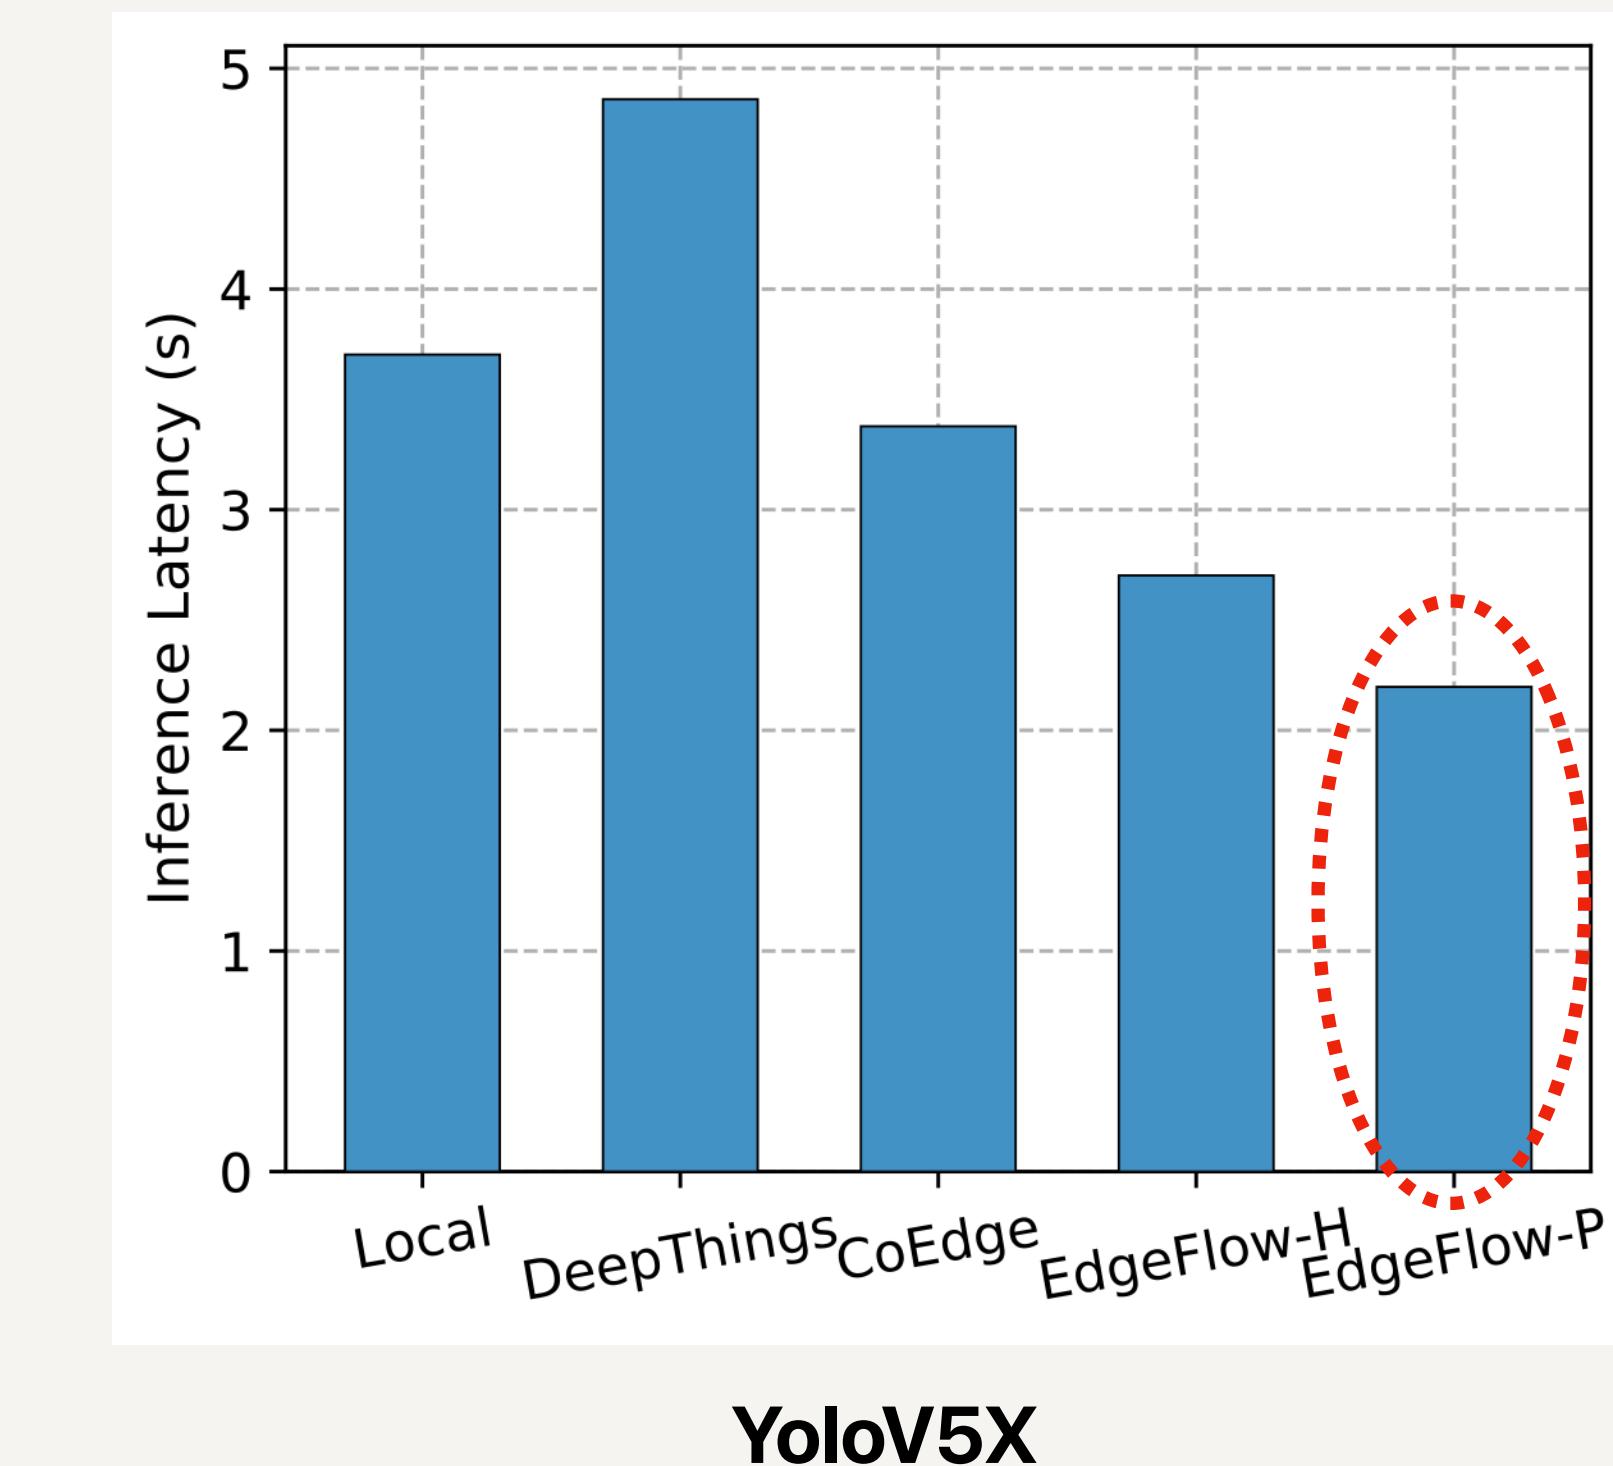
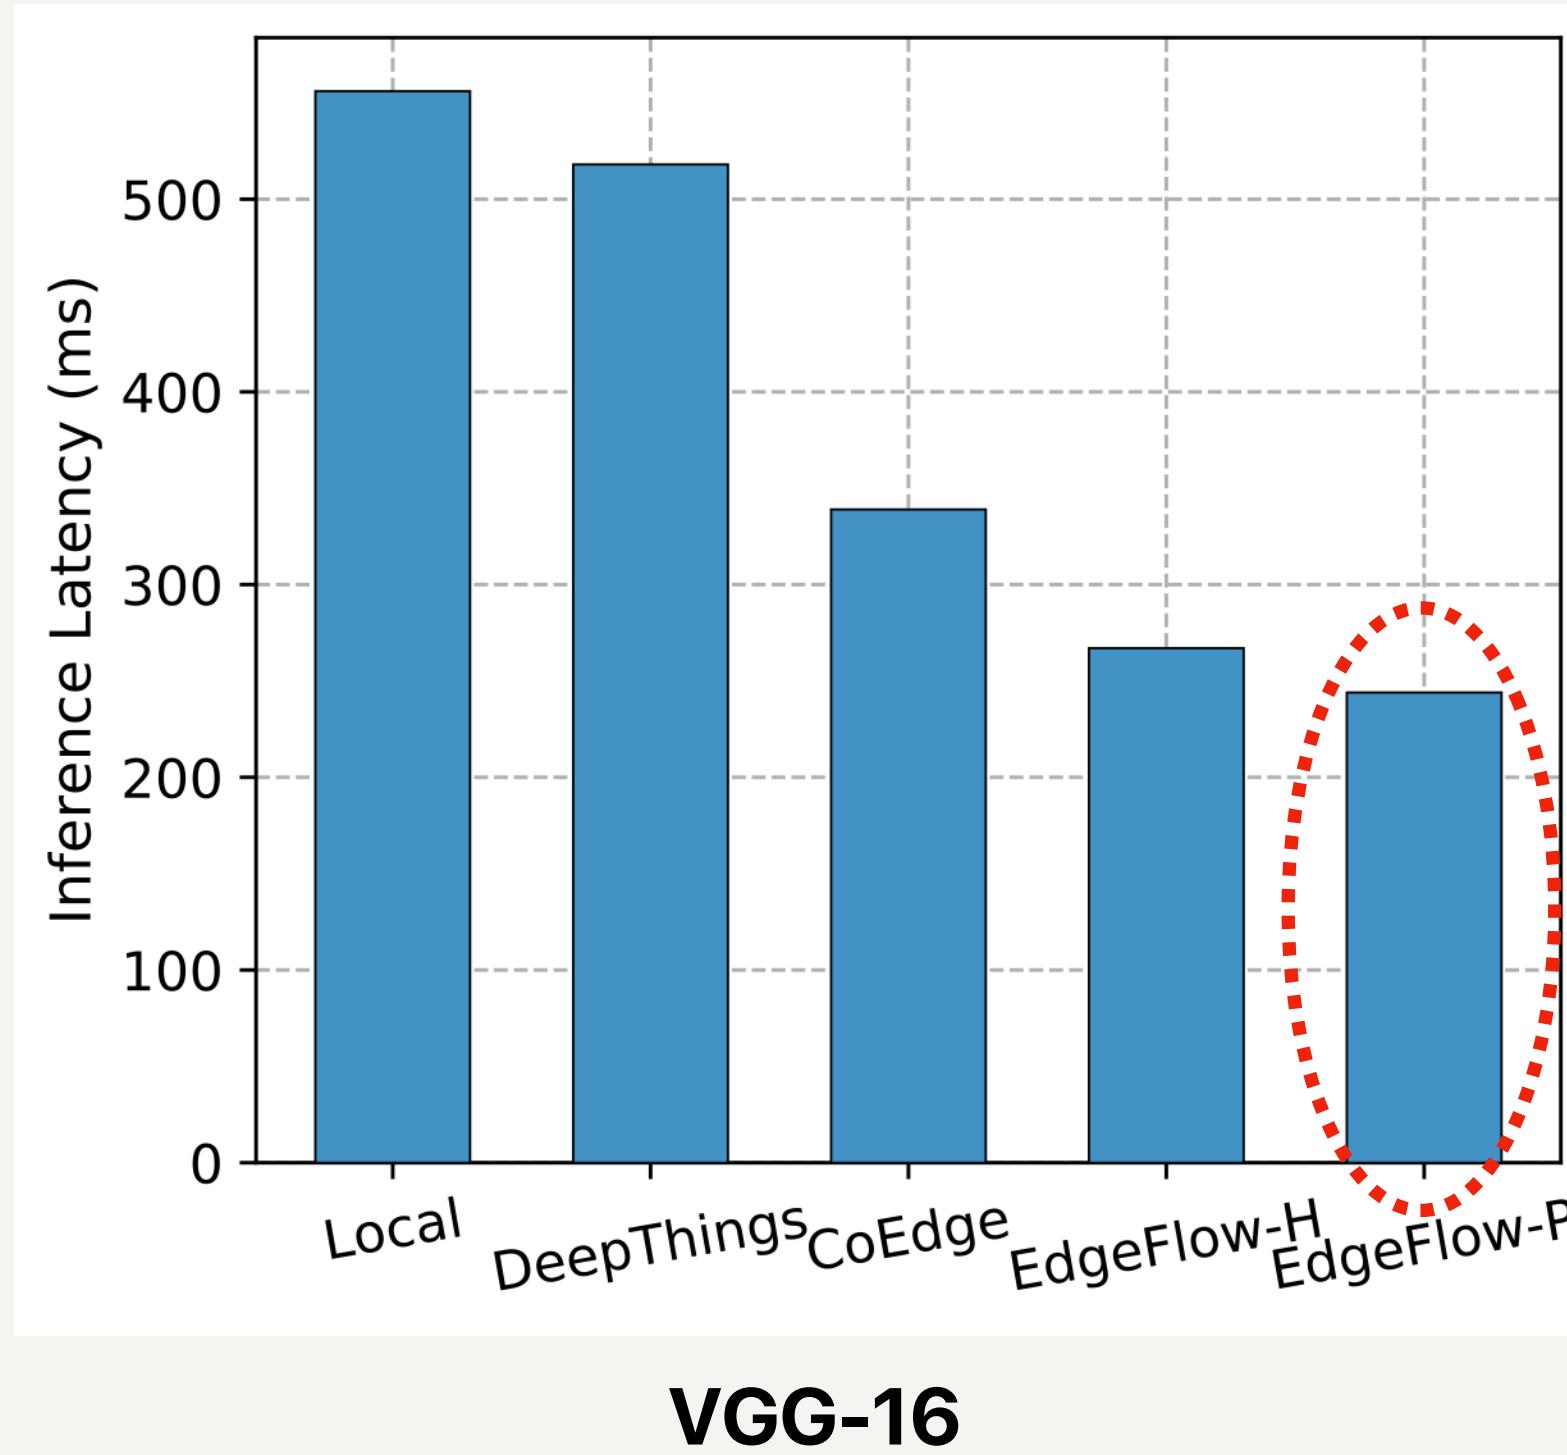


# Evaluation

- ▶ 2 deep learning models
  - ▶ VGG-16: Classic image classification model in chain structure
  - ▶ YoloV5X: Latest object detection model with complicated structure
- ▶ 6 heterogeneous virtual machines
- ▶ Baselines
  - ▶ Local: deploy the model on a single device
  - ▶ Existing methods: DeepThings and CoEdge

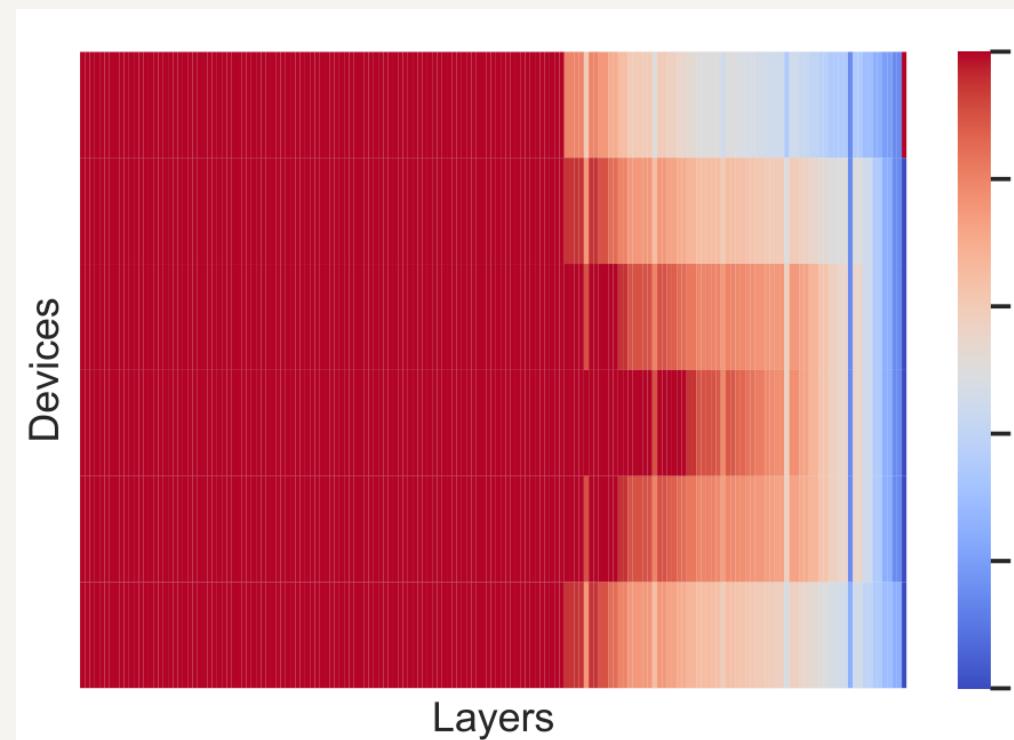
# Evaluation

Proposed method (EdgeFlow-P) achieves lowest inference latency with both models

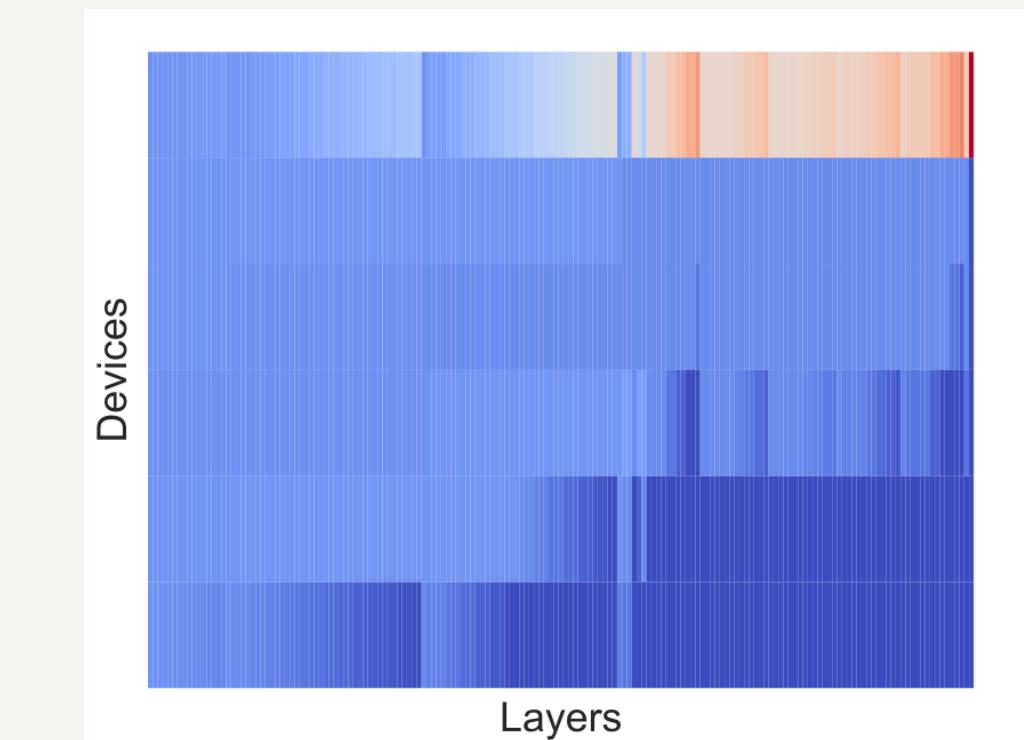


# Evaluation

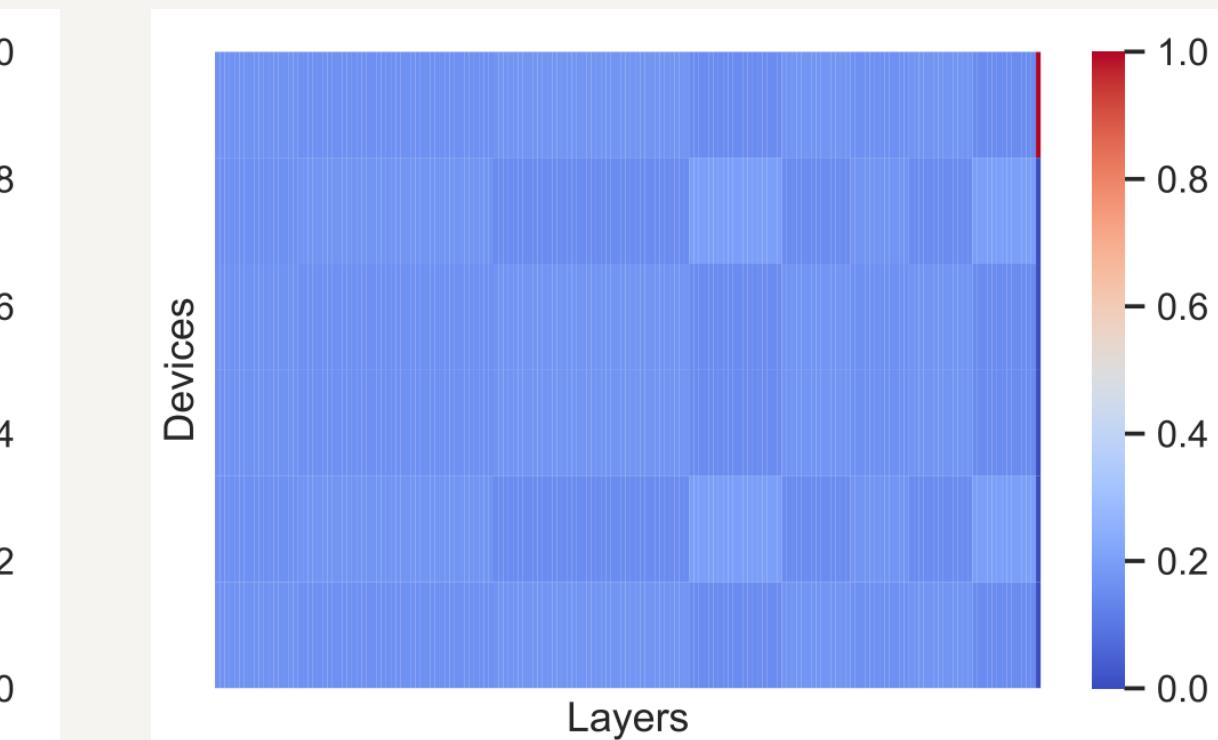
- ▶ Partition scheme of YoloV5
  - ▶ DeepThings: redundant computation in the early layers
  - ▶ CoEdge: workload gradually concentrates on a single device
  - ▶ EdgeFlow: relatively even distribution among devices



**DeepThings**



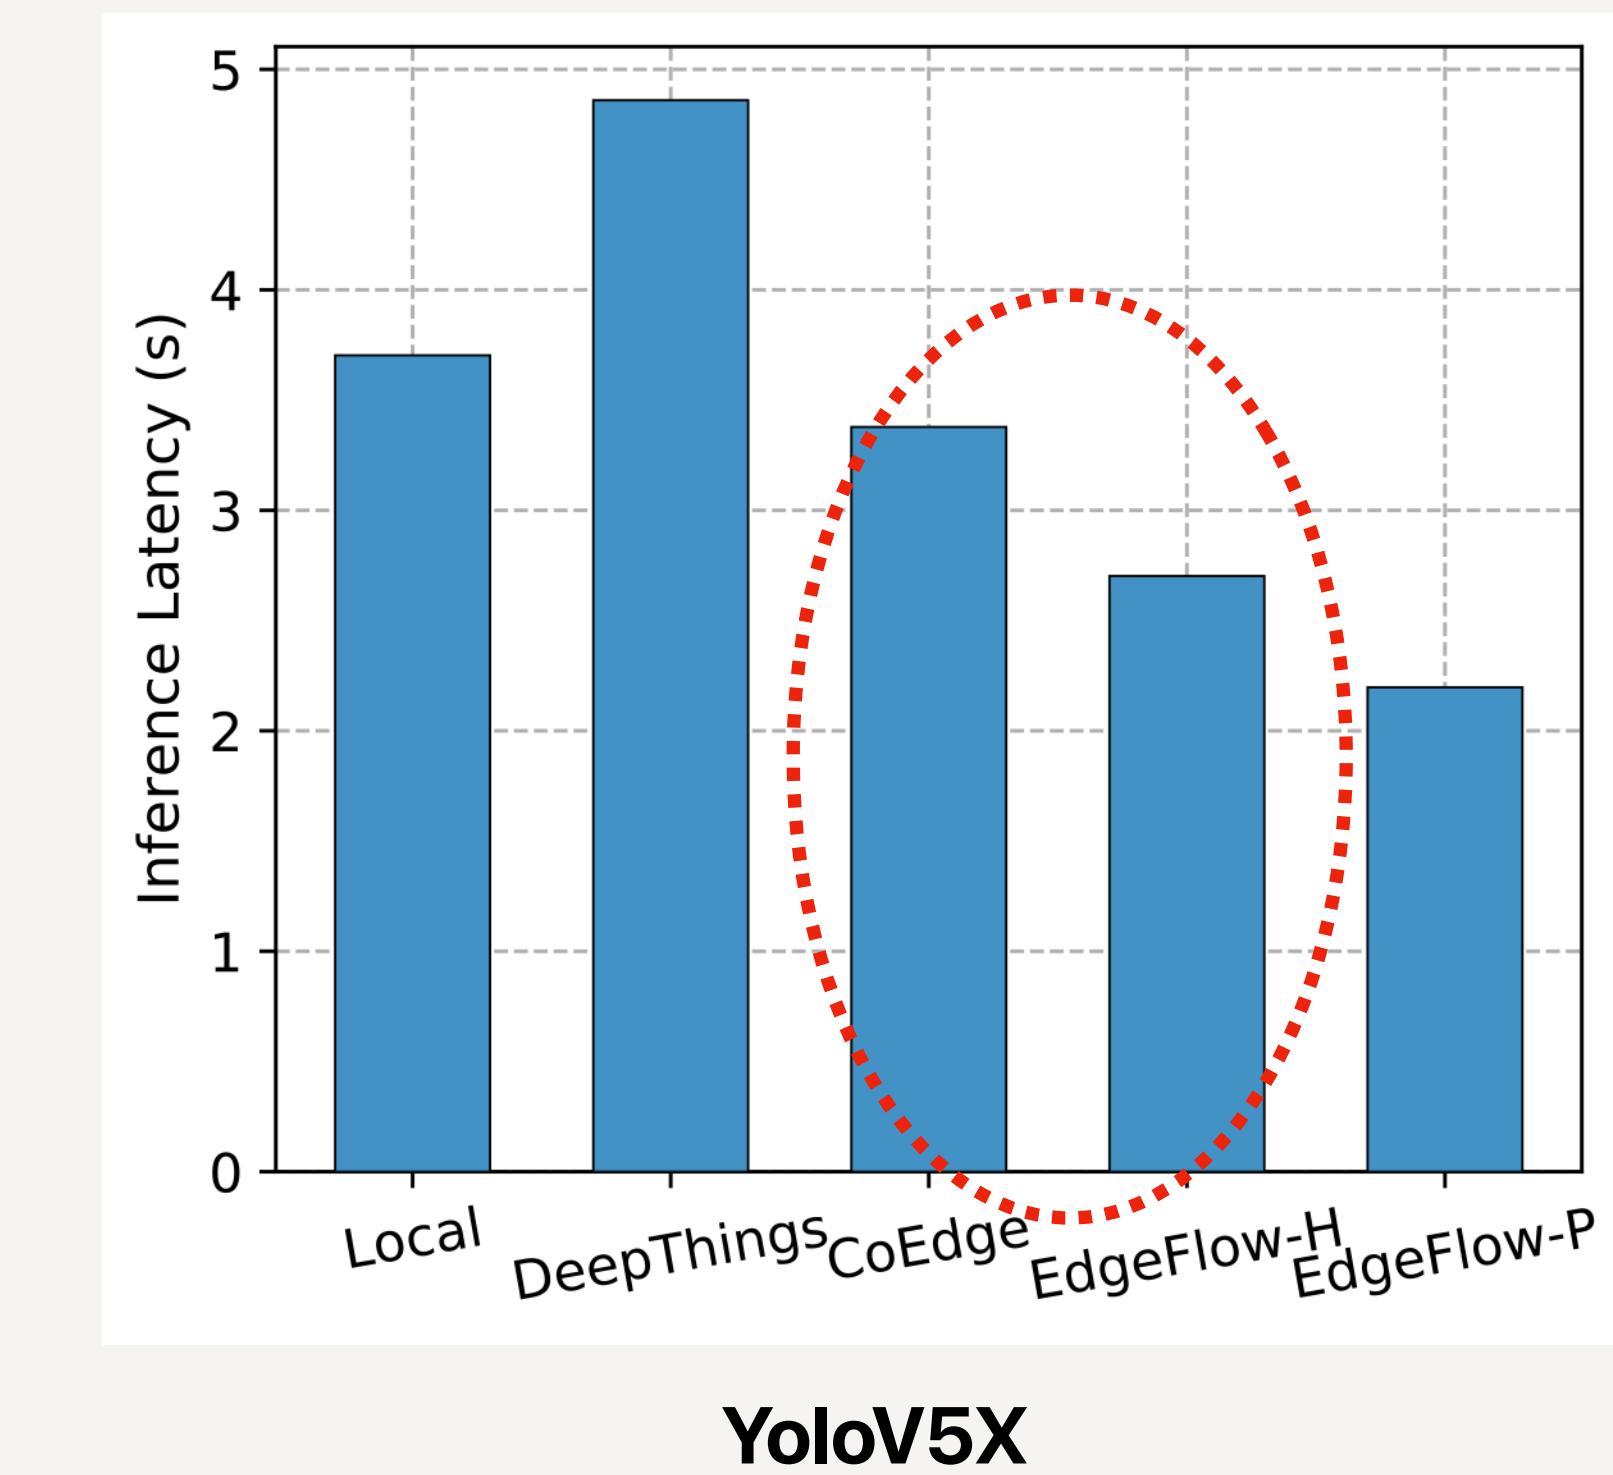
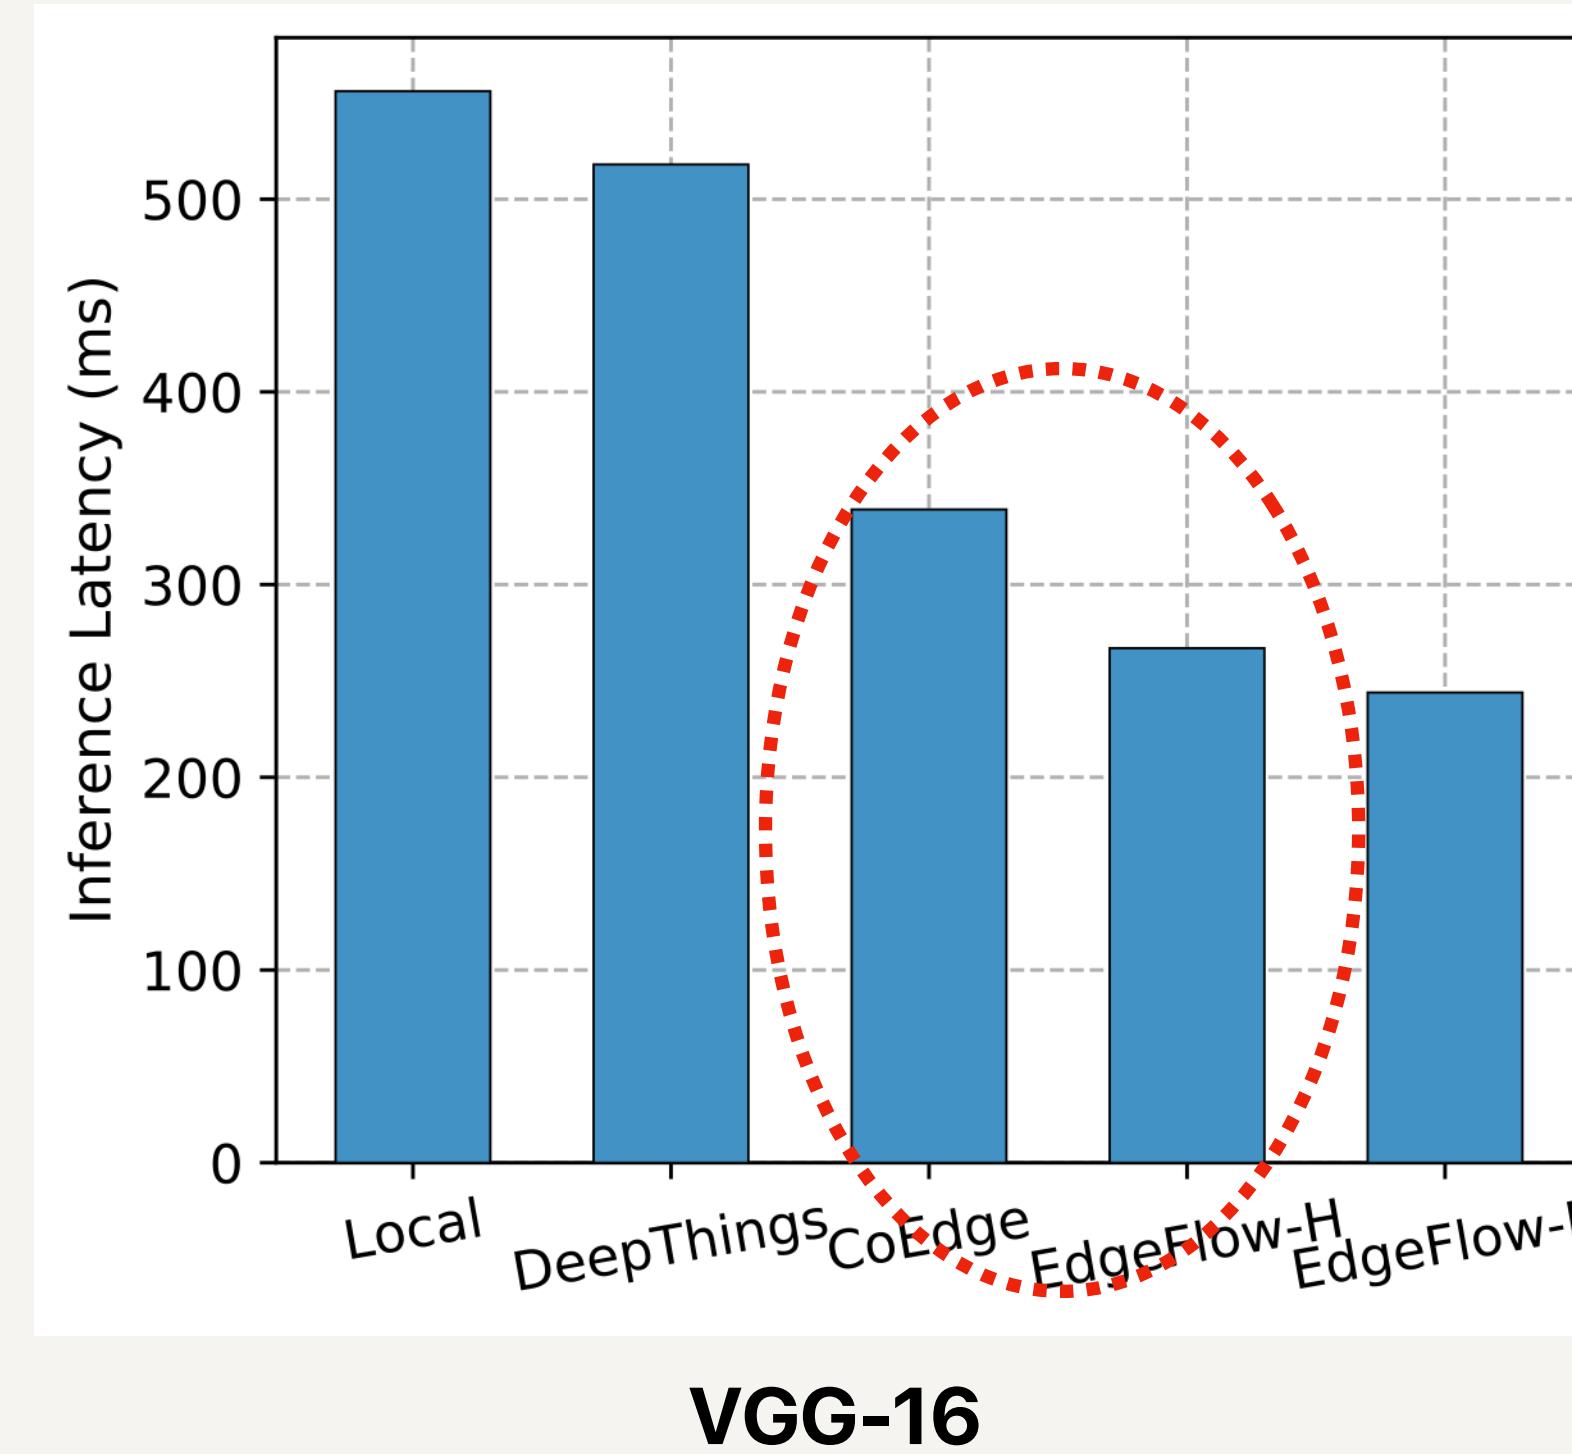
**CoEdge**



**EdgeFlow**

# Evaluation

EdgeFlow-H and CoEdge share the same partition scheme, yet still faster than CoEdge



# Conclusion

- ▶ The model structure significantly affects the performance of existing distributed inference systems.
- ▶ *EdgeFlow* breaks the layer into execution units, and maintain the complicated layer dependencies by controlling the flow of intermediate results.
- ▶ Evaluation results show *EdgeFlow* has a distinct advantage, especially with complicated DAG-structured model

Contact: [ch.hu@mail.utoronto.ca](mailto:ch.hu@mail.utoronto.ca)